

```
!pip install bayesian-optimization
```

```
Collecting bayesian-optimization
  Downloading bayesian_optimization-1.4.3-py3-none-any.whl (18 kB)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.25.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.11.4)
Requirement already satisfied: scikit-learn>=0.18.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization)
Collecting colorama>=0.4.6 (from bayesian-optimization)
  Downloading colorama-0.4.6-py2.py3-none-any.whl (25 kB)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (3.2.0)
Installing collected packages: colorama, bayesian-optimization
Successfully installed bayesian-optimization-1.4.3 colorama-0.4.6
```

```
!git clone https://github.com/808ss/thesis.git
```

```
Cloning into 'thesis'...
remote: Enumerating objects: 27, done.
remote: Counting objects: 100% (27/27), done.
remote: Compressing objects: 100% (26/26), done.
remote: Total 27 (delta 0), reused 0 (delta 0), pack-reused 0
Receiving objects: 100% (27/27), 311.32 KiB | 1.26 MiB/s, done.
```

```
import numpy as np
import pandas as pd
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import seaborn as sns
import itertools
from bayes_opt import BayesianOptimization
```

```
random_seed = 808
np.random.seed(random_seed)
```

✓ MBBR

✓ Importing MBBR and Splitting

```
MBBR = pd.read_csv('thesis/MBBR-Chlorination.csv')
MBBR.drop(columns='Date', inplace=True)
```

```
X_orig_MBBR = MBBR.drop(columns='Residual chlorine\n(ppm)')
y_orig_MBBR = MBBR['Residual chlorine\n(ppm)']
X_train_orig_MBBR, X_test_orig_MBBR, y_train_orig_MBBR, y_test_orig_MBBR = train_test_split(X_orig_MBBR,
                                                                                          y_orig_MBBR,
                                                                                          test_size = 0.3,
                                                                                          random_state=808)
```

```
df_train_orig_MBBR = pd.concat([X_train_orig_MBBR, y_train_orig_MBBR], axis=1)
df_test_orig_MBBR = pd.concat([X_test_orig_MBBR, y_test_orig_MBBR], axis=1)
```

✓ Data Analysis for Raw Dataset

```
missing_rate_MBBR = [(MBBR.isnull().sum()[val]/MBBR.shape[0])*100 for val in range(0, MBBR.shape[1])]
```

```
pd.options.display.float_format = '{:,.2f}'.format
MBBR_transposed = MBBR.describe().T
MBBR_transposed['Missingness Rate'] = missing_rate_MBBR
```

```
MBBR_transposed
```



	count	mean	std	min	25%	50%	75%	max	Missingness Rate
Flow Rate Influent (m3/d)	332.00	4,787.53	2,211.48	197.00	3,344.00	4,709.50	6,232.00	11,147.00	0.0%
Total Coliform Influent (MPN/100mL)	270.00	290,896,939.26	733,941,441.61	1,600.00	17,250,000.00	40,500,000.00	160,000,000.00	5,200,000,000.00	18.6%
Total Coliform Effluent (MPN/100mL)	329.00	733,375.06	9,402,984.69	0.00	2.00	10.00	471.00	143,900,000.00	0.9%
Fecal Coliform Influent (MPN/100mL)	103.00	236,377,087.38	621,705,589.64	230,000.00	8,550,000.00	23,000,000.00	37,650,000.00	3,000,000,000.00	68.9%
Fecal Coliform Effluent (MPN/100mL)	171.00	746.87	3,947.85	2.00	10.00	10.00	10.00	24,196.00	48.4%
BOD Influent (ppm)	273.00	152.40	148.02	8.00	68.00	119.00	196.00	1,425.00	17.7%
BOD Pre-chlorination\n(ppm)	274.00	11.28	12.82	1.00	4.00	8.00	14.00	119.00	17.4%

▼ Data Analysis for Training Set (Pre-Imputation)

```
missing_rate_train_orig_MBBR = [(df_train_orig_MBBR.isnull().sum()[val]/df_train_orig_MBBR.shape[0])*100 for val in range(0,df_train_orig_MBBR.shape[1])]

pd.options.display.float_format = '{:,.2f}'.format
#pd.set_option('display.float_format', '{:e}'.format)
df_train_orig_MBBR_transposed = df_train_orig_MBBR.describe().T
df_train_orig_MBBR_transposed['Missingness Rate'] = missing_rate_train_orig_MBBR

df_train_orig_MBBR_transposed
```



	count	mean	std	min	25%	50%	75%	max	Missingness Rate
Flow Rate Influent (m3/d)	232.00	4,882.74	2,204.80	197.00	3,344.00	4,762.00	6,349.00	10,999.00	0.0%
Total Coliform Influent (MPN/100mL)	185.00	315,522,010.81	790,769,090.66	16,000.00	18,000,000.00	41,000,000.00	160,000,000.00	5,200,000,000.00	20.2%
Total Coliform Effluent (MPN/100mL)	230.00	1,045,242.88	11,238,969.87	0.00	2.25	10.00	1,280.75	143,900,000.00	0.8%
Fecal Coliform Influent (MPN/100mL)	68.00	298,124,117.65	669,480,560.27	230,000.00	10,400,000.00	24,000,000.00	40,950,000.00	2,600,000,000.00	70.6%
Fecal Coliform Effluent (MPN/100mL)	120.00	892.02	4,352.93	2.00	10.00	10.00	10.00	24,196.00	48.2%
BOD Influent (ppm)	187.00	162.30	167.60	8.00	70.50	122.00	199.00	1,425.00	19.4%
BOD Pre-chlorination\n(ppm)	188.00	11.12	12.31	1.00	5.00	8.00	14.00	119.00	18.9%

▼ Data Analysis for Testing Set (Pre-imputation)

```
missing_rate_test_orig_MBBR = [(df_test_orig_MBBR.isnull().sum()[val]/df_test_orig_MBBR.shape[0])*100 for val in range(0,df_test_orig_MBBR.shape[1])]

#pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
df_test_orig_MBBR_transposed = df_test_orig_MBBR.describe().T
df_test_orig_MBBR_transposed['Missingness Rate'] = missing_rate_test_orig_MBBR

df_test_orig_MBBR_transposed
```



	count	mean	std	min	25%	50%	75%	max	Missing
Flow Rate Influent (m3/d)	1.000000e+02	4.566650e+03	2.222244e+03	2.170000e+02	3.343750e+03	4.641500e+03	6.112250e+03	1.114700e+04	0.000000
Total Coliform Influent (MPN/100mL)	8.500000e+01	2.373012e+08	5.924907e+08	1.600000e+03	1.700000e+07	4.000000e+07	1.600000e+08	3.500000e+09	1.500000
Total Coliform Effluent (MPN/100mL)	9.900000e+01	8.833667e+03	3.764360e+04	0.000000e+00	2.000000e+00	1.000000e+01	9.000000e+01	2.419600e+05	1.000000
Fecal Coliform Influent (MPN/100mL)	3.500000e+01	1.164114e+08	5.038721e+08	1.000000e+06	7.450000e+06	1.700000e+07	3.045000e+07	3.000000e+09	6.500000
Fecal Coliform Effluent (MPN/100mL)	5.100000e+01	4.053529e+02	2.779390e+03	2.000000e+00	2.000000e+00	1.000000e+01	1.000000e+01	1.986300e+04	4.900000
BOD Influent (ppm)	8.600000e+01	1.308605e+02	8.921358e+01	1.900000e+01	6.600000e+01	1.060000e+02	1.830000e+02	4.090000e+02	1.400000
BOD Pre-chlorination\n(ppm)	8.600000e+01	1.162791e+01	1.393434e+01	1.000000e+00	4.000000e+00	8.000000e+00	1.400000e+01	1.080000e+02	1.400000

▼ Data Imputation

▼ Exporting Datasets to R

```
df_train_orig_MBBR.to_csv('MBBR_train_set.csv',index=False)
df_test_orig_MBBR.to_csv('MBBR_test_set.csv',index=False)
```

Export to R for mixgb

▼ Mixgb imputation

```
1 library(mixgb)
2 library(openxlsx)
3 set.seed(808)
4
5 MBBR_train_set <- read.csv("C:/Users/nikko/PycharmProjects/Thesis/MBBR_train_set.csv")
6 MBBR_test_set <- read.csv("C:/Users/nikko/PycharmProjects/Thesis/MBBR_test_set.csv")
7
8 MBBR_train_set_df = as.data.frame(MBBR_train_set)
9 MBBR_test_set_df = as.data.frame(MBBR_test_set)
10
11 clean_MBBR_train_set_df <- data_clean(MBBR_train_set_df)
12 clean_MBBR_test_set_df <- data_clean(MBBR_test_set_df)
13
14 cv.results_1 <- mixgb_cv(data = clean_MBBR_train_set_df, nrounds = 5000, verbose = FALSE)
15 cv.results_1$evaluation.log
16 cv.results_1$best.nrounds
17
18 mixgb_obj <- mixgb(data = clean_MBBR_train_set_df, m = 5, nrounds = cv.results_1$best.nrounds, save.models = TRUE)
19 MBBR_train_imputed <- mixgb_obj$imputed.data
20
21 MBBR_test_imputed <- impute_new(object = mixgb_obj, newdata = clean_MBBR_test_set_df)
22
23 write.xlsx(MBBR_train_imputed[[1]], file = 'mbbr_m1_imputed_train.xlsx')
24 write.xlsx(MBBR_train_imputed[[2]], file = 'mbbr_m2_imputed_train.xlsx')
25 write.xlsx(MBBR_train_imputed[[3]], file = 'mbbr_m3_imputed_train.xlsx')
26 write.xlsx(MBBR_train_imputed[[4]], file = 'mbbr_m4_imputed_train.xlsx')
27 write.xlsx(MBBR_train_imputed[[5]], file = 'mbbr_m5_imputed_train.xlsx')
28
29 write.xlsx(MBBR_test_imputed[[1]], file = 'mbbr_m1_imputed_test.xlsx')
30 write.xlsx(MBBR_test_imputed[[2]], file = 'mbbr_m2_imputed_test.xlsx')
31 write.xlsx(MBBR_test_imputed[[3]], file = 'mbbr_m3_imputed_test.xlsx')
32 write.xlsx(MBBR_test_imputed[[4]], file = 'mbbr_m4_imputed_test.xlsx')
33 write.xlsx(MBBR_test_imputed[[5]], file = 'mbbr_m5_imputed_test.xlsx')
```

▼ Import imputed datasets from R

```
dfs = []
for val in range(1,6):
    source = f'thesis/mbbr_m{val}_imputed_train.xlsx'
```

```
dfs.append(pd.read_excel(source))

average_MBBR_train = pd.concat(dfs).groupby(level=0).mean()

dfs = []
for val in range(1,6):
    source = f'thesis/mbbr_m{val}_imputed_test.xlsx'
    dfs.append(pd.read_excel(source))

average_MBBR_test = pd.concat(dfs).groupby(level=0).mean()
```

▼ Data Analysis for Training Set (Post-Imputation)

```
#pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
average_MBBR_train_transposed = average_MBBR_train.describe().T
```

average_MBBR_train_transposed

	count	mean	std	min	25%	50%	75%	
Flow.Rate.Influent..m3.d.	2.320000e+02	4.882741e+03	2.204801e+03	1.970000e+02	3.344000e+03	4.762000e+03	6.349000e+03	1.09990
Total.Coliform.Influent..MPN.100mL.	2.320000e+02	3.480306e+08	7.784700e+08	1.600000e+04	2.175000e+07	5.400000e+07	2.200000e+08	5.20000
Total.Coliform.Effluent..MPN.100mL.	2.320000e+02	1.036336e+06	1.119062e+07	0.000000e+00	2.750000e+00	1.000000e+01	1.600000e+03	1.43900
Fecal.Coliform.Influent..MPN.100mL.	2.320000e+02	1.965840e+08	5.001635e+08	2.300000e+05	1.428500e+07	2.740000e+07	3.821000e+07	2.60000
Fecal.Coliform.Effluent..MPN.100mL.	2.320000e+02	4.182178e+03	8.256544e+03	2.000000e+00	8.800000e+00	1.000000e+01	2.282500e+02	2.41960
BOD.Influent..ppm.	2.320000e+02	1.585655e+02	1.533443e+02	8.000000e+00	7.585000e+01	1.250000e+02	1.950500e+02	1.42500
BOD.Pre.chlorination..ppm.	2.320000e+02	1.199138e+01	1.246310e+01	1.000000e+00	5.000000e+00	9.000000e+00	1.500000e+01	1.19000
COD.Influent..ppm.	2.320000e+02	3.458931e+02	5.439009e+02	1.300000e+01	1.750000e+02	2.510000e+02	3.745000e+02	7.73400
COD.Pre.chlorination..ppm.	2.320000e+02	4.940431e+01	3.795535e+01	5.000000e+00	2.475000e+01	4.150000e+01	6.300000e+01	3.43000
TSS.Pre.chlorination..ppm.	2.320000e+02	1.756897e+01	2.094736e+01	1.000000e+00	6.000000e+00	1.200000e+01	2.000000e+01	1.60000
pH.Pre.chlorination	2.320000e+02	7.185121e+00	3.010609e-01	6.120000e+00	7.000000e+00	7.200000e+00	7.362500e+00	8.38000
Chlorine.dosage..L.d.	2.320000e+02	8.721017e+02	4.852711e+02	0.000000e+00	6.000000e+02	8.500000e+02	1.160000e+03	2.80000
Residual.chlorine..ppm.	2.320000e+02	2.082141e+00	1.913163e+00	0.000000e+00	3.815000e-01	1.205700e+00	3.917150e+00	5.48000

▼ Data Analysis for Testing Set (Post-Imputation)

```
pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
average_MBBR_test_transposed = average_MBBR_test.describe().T
```

average_MBBR_test_transposed



	count	mean	std	min	25%	50%	75%	
Flow.Rate.Influent..m3.d.	100.00	4,566.65	2,222.24	217.00	3,343.75	4,641.50	6,112.25	11,147
Total.Coliiform.Influent..MPN.100mL.	100.00	288,973,956.00	647,069,268.14	1,600.00	22,750,000.00	45,000,000.00	200,000,000.00	3,500,000,000
Total.Coliiform.Effluent..MPN.100mL.	100.00	8,933.07	37,466.19	0.00	2.00	10.00	134.75	241,960
Fecal.Coliiform.Influent..MPN.100mL.	100.00	147,284,740.00	427,084,482.81	1,000,000.00	12,000,000.00	25,320,000.00	41,560,000.00	3,000,000,000
Fecal.Coliiform.Effluent..MPN.100mL.	100.00	2,363.05	6,219.15	2.00	10.00	10.00	107.30	24,196
BOD.Influent..ppm.	100.00	134.13	88.52	19.00	66.45	108.40	188.25	408
BOD.Pre.chlorination..ppm.	100.00	12.06	13.39	1.00	5.00	9.00	15.00	108
COD.Influent..ppm.	100.00	268.66	192.32	41.00	135.75	210.00	357.25	1,256
COD.Pre.chlorination..ppm.	100.00	45.03	47.09	5.00	20.00	29.70	53.55	314
TSS.Pre.chlorination..ppm.	100.00	18.00	24.39	1.00	5.00	10.00	19.25	164
pH.Pre.chlorination	100.00	7.21	0.37	5.28	7.10	7.21	7.40	8
Chlorine.dosage..L.d.	100.00	809.32	498.65	0.00	488.70	748.90	1,005.05	2,900
Residual.chlorine..ppm.	100.00	2.14	1.69	0.01	0.49	2.05	3.57	5

✎ Exhaustive Feature Selection

✎ For Imputed Dataset

```
pd.reset_option('display.float_format')
```

```
X_train_MBBR = average_MBBR_train.drop(columns='Residual.chlorine..ppm.')
y_train_MBBR = average_MBBR_train['Residual.chlorine..ppm.']
X_test_MBBR = average_MBBR_test.drop(columns='Residual.chlorine..ppm.')
y_test_MBBR = average_MBBR_test['Residual.chlorine..ppm.']}
```

```
features_wo_chlorine_dosage = X_train_MBBR.columns[:-1]
features_wo_chlorine_dosage
```

```
Index(['Flow.Rate.Influent..m3.d.', 'Total.Coliiform.Influent..MPN.100mL.',
      'Total.Coliiform.Effluent..MPN.100mL.',
      'Fecal.Coliiform.Influent..MPN.100mL.',
      'Fecal.Coliiform.Effluent..MPN.100mL.', 'BOD.Influent..ppm.',
      'BOD.Pre.chlorination..ppm.', 'COD.Influent..ppm.',
      'COD.Pre.chlorination..ppm.', 'TSS.Pre.chlorination..ppm.',
      'pH.Pre.chlorination'],
      dtype='object')
```

```
# Generate all combinations of the other features
combinations = []
for r in range(1, len(features_wo_chlorine_dosage) + 1):
    combinations.extend(itertools.combinations(features_wo_chlorine_dosage, r))
```

```
# Add the first feature to each combination
combinations = [(X_train_MBBR.columns[-1],) + combo for combo in combinations]
```

```
params = {'objective': 'reg:squarederror'}
```

```
results = []
for combo in combinations:
    dtrain = xgb.DMatrix(X_train_MBBR[list(combo)], label=y_train_MBBR)
    cv_result = xgb.cv(params, dtrain, num_boost_round=10, nfold=5, metrics='rmse', seed=808)
    last_round_metrics = cv_result.iloc[-1]
    results.append([combo, last_round_metrics['train-rmse-mean'], last_round_metrics['test-rmse-mean'],
                  last_round_metrics['train-rmse-std'], last_round_metrics['test-rmse-std']])
```

```
results_df_MBBR = pd.DataFrame(results, columns=['Combination', 'Train RMSE', 'Validation RMSE', 'Train RMSE Std. Dev.', 'Valid
```

```
results_df_MBBR.sort_values(by='Validation RMSE')
```



	Combination	Train RMSE	Validation RMSE	Train RMSE Std. Dev.	Validation RMSE Std. Dev
1851	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.348480	1.394428	0.028309	0.143395
1167	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.315027	1.411094	0.044549	0.073128
1636	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.302548	1.415478	0.028380	0.069078
1502	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.344748	1.417683	0.018418	0.113160
1493	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.340146	1.418291	0.035193	0.142621
...
193	(Chlorine.dosage..L.d., Fecal.Coliiform.Influen...	0.788887	2.085864	0.043836	0.179699
187	(Chlorine.dosage..L.d., Fecal.Coliiform.Influen...	0.742502	2.099164	0.053870	0.108392
3	(Chlorine.dosage..L.d., Fecal.Coliiform.Influen...	0.966307	2.111154	0.030007	0.200713
8	(Chlorine.dosage..L.d., COD.Pre.chlorination....	1.070077	2.117816	0.101896	0.117162
42	(Chlorine.dosage..L.d., Fecal.Coliiform.Influen...	0.857085	2.160103	0.053470	0.113003

2047 rows x 5 columns

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[0:3]
```



	Combination	Train RMSE	Validation RMSE	Train RMSE Std. Dev.	Validation RMSE Std. Dev
1851	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.348480	1.394428	0.028309	0.143395
1167	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.315027	1.411094	0.044549	0.073128
1636	(Chlorine.dosage..L.d., Flow.Rate.Influent..m3...	0.302548	1.415478	0.028380	0.069078

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[0]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'Total.Coliiform.Effluent..MPN.100mL.',
 'Fecal.Coliiform.Effluent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'BOD.Pre.chlorination..ppm.',
 'COD.Influent..ppm.',
 'TSS.Pre.chlorination..ppm.')
```

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[1]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Effluent..MPN.100mL.',
 'Fecal.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'BOD.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[2]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Effluent..MPN.100mL.',
 'Fecal.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'BOD.Pre.chlorination..ppm.',
 'TSS.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
optimal_features_MBBR = results_df_MBBR.sort_values(by='Validation RMSE').iloc[0]['Combination']
optimal_features_MBBR
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'Total.Coliiform.Effluent..MPN.100mL.',
 'Fecal.Coliiform.Effluent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'BOD.Pre.chlorination..ppm.',
```

```
'COD.Influent..ppm.',
'TSS.Pre.chlorination..ppm.')

results_df_MBBR['count'] = results_df_MBBR['Combination'].apply(lambda x: len(x))
results_df_MBBR.to_csv('MBBR Exhaustive Feature Selection.csv', index=False)
```

For Raw Dataset

```
non_imputed_mask_MBBR_train = ~np.isnan(y_train_orig_MBBR)
non_imputed_mask_MBBR_test = ~np.isnan(y_test_orig_MBBR)

X_train_MBBR_dropped = X_train_orig_MBBR[non_imputed_mask_MBBR_train]
y_train_MBBR_dropped = y_train_orig_MBBR[non_imputed_mask_MBBR_train]
X_test_MBBR_dropped = X_test_orig_MBBR[non_imputed_mask_MBBR_test]
y_test_MBBR_dropped = y_test_orig_MBBR[non_imputed_mask_MBBR_test]

features_wo_chlorine_dosage_dropped = X_train_MBBR_dropped.columns[:-1]
features_wo_chlorine_dosage_dropped

Index(['Flow Rate Influent (m3/d)', 'Total Coliform Influent (MPN/100mL)',
      'Total Coliform Effluent (MPN/100mL)',
      'Fecal Coliform Influent (MPN/100mL)',
      'Fecal Coliform Effluent (MPN/100mL)', 'BOD Influent (ppm)',
      'BOD Pre-chlorination\n(ppm)', 'COD Influent (ppm)',
      'COD Pre-chlorination\n(ppm)', 'TSS Pre-chlorination (ppm)',
      'pH Pre-chlorination'],
      dtype='object')

# Generate all combinations of the other features
combinations = []
for r in range(1, len(features_wo_chlorine_dosage_dropped) + 1):
    combinations.extend(itertools.combinations(features_wo_chlorine_dosage_dropped, r))

# Add the first feature to each combination
combinations = [(X_train_MBBR_dropped.columns[-1],) + combo for combo in combinations]

params = {'objective': 'reg:squarederror'}

results = []
for combo in combinations:
    dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(combo)], label=y_train_MBBR_dropped)
    cv_result = xgb.cv(params, dtrain, num_boost_round=10, nfold=5, metrics='rmse', seed=808)
    last_round_metrics = cv_result.iloc[-1]
    results.append([combo, last_round_metrics['train-rmse-mean'], last_round_metrics['test-rmse-mean'],
                  last_round_metrics['train-rmse-std'], last_round_metrics['test-rmse-std']])

results_df_MBBR_dropped = pd.DataFrame(results, columns=['Combination', 'Train RMSE', 'Validation RMSE', 'Train RMSE Std. Dev.',
                                                       'Validation RMSE Std. Dev.'])

results_df_MBBR_dropped.sort_values(by='Validation RMSE')
```

↗

	Combination	Train RMSE	Validation RMSE	Train RMSE Std. Dev.	Validation RMSE Std. Dev
31	(Chlorine dosage (L/d), Total Coliform Effluen...	1.022659	1.512094	0.031956	0.180507
2	(Chlorine dosage (L/d), Total Coliform Effluen...	1.031525	1.520116	0.038504	0.184722
32	(Chlorine dosage (L/d), Total Coliform Effluen...	0.636437	1.554748	0.050116	0.192632
160	(Chlorine dosage (L/d), Total Coliform Effluen...	0.612772	1.574227	0.047039	0.187918
574	(Chlorine dosage (L/d), Flow Rate Influent (m3...	0.424320	1.577279	0.044534	0.284987
...
1000	(Chlorine dosage (L/d), Fecal Coliform Influen...	0.420868	2.110952	0.038620	0.188357
139	(Chlorine dosage (L/d), Total Coliform Influen...	0.576095	2.111329	0.039815	0.206058
44	(Chlorine dosage (L/d), Fecal Coliform Influen...	0.804355	2.112326	0.033380	0.167510
525	(Chlorine dosage (L/d), Fecal Coliform Influen...	0.415596	2.112704	0.030191	0.144628
58	(Chlorine dosage (L/d), BOD Pre-chlorination\n...	0.648844	2.153476	0.061506	0.111363

2047 rows x 5 columns

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0:3]
```

	Combination	Train RMSE	Validation RMSE	Train RMSE Std. Dev.	Validation RMSE Std. Dev
31	(Chlorine dosage (L/d), Total Coliform Effluen...	1.022659	1.512094	0.031956	0.180507
2	(Chlorine dosage (L/d), Total Coliform Effluen...	1.031525	1.520116	0.038504	0.184722
32	(Chlorine dosage (L/d), Total Coliform Effluen...	0.636437	1.554748	0.050116	0.192632

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0]['Combination']
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'Fecal Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[1]['Combination']
```

```
('Chlorine dosage (L/d)', 'Total Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[2]['Combination']
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'BOD Influent (ppm)')
```

```
optimal_features_MBBR_dropped = results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0]['Combination']
optimal_features_MBBR_dropped
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'Fecal Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped['count'] = results_df_MBBR_dropped['Combination'].apply(lambda x: len(x))
results_df_MBBR_dropped.to_csv('MBBR Dropped Exhaustive Feature Selection.csv', index=False)
```

✓ Hyperparameter Optimization

✓ For Imputed Dataset

```
# Convert the data into DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR[list(optimal_features_MBBR)], label=y_train_MBBR)

# Define the function to be optimized
def xgb_evaluate(eta, alpha, lambd, gamma, subsample, col_subsample, max_depth):
    eta = 10**eta
    alpha = 10**alpha
    lambd = 10**lambd
    gamma = 10**gamma
    max_depth = int(round(2**max_depth))

    params = {'eval_metric': 'rmse',
              'objective': 'reg:squarederror',
              'max_depth': max_depth,
              'eta': eta,
              'gamma': gamma,
              'subsample': subsample,
              'alpha': alpha,
              'lambda': lambd,
              'colsample_bytree': col_subsample,}

    cv_result = xgb.cv(params, dtrain, num_boost_round=1000, nfold=5, early_stopping_rounds=30, seed=808)
    return -1.0 * cv_result['test-rmse-mean'].iloc[-1]

# Specify the hyperparameters to be tuned
xgb_bo_MBBR = BayesianOptimization(xgb_evaluate, {'eta': (-3, 0),
                                                  'alpha': (-6, 0.3),
                                                  'lambd': (-6, 0.3),
                                                  'gamma': (-6, 1.8),
                                                  'subsample': (0.5, 1),
                                                  'col_subsample': (0.3, 1),
```



```
'max_depth': (1, 3)},
random_state=808)
```

```
# Optimize the hyperparameters
```

```
xgb_bo_MBBR.maximize(n_iter=1000, init_points=10)# Convert the data into DMatrix format
```

iter	target	alpha	col_su...	eta	gamma	lambda	max_depth	subsample
1	-1.402	0.04075	0.4513	-2.68	-1.662	-1.582	2.026	0.7673
2	-1.348	-4.514	0.7529	-1.843	-2.2	-1.339	1.596	0.5436
3	-1.412	-1.108	0.5069	-1.136	-4.974	-0.5216	2.693	0.8202
4	-1.522	-3.147	0.6275	-2.063	1.604	-0.4504	1.466	0.7294
5	-1.52	-2.356	0.4052	-0.3806	-0.09483	-5.01	1.643	0.6674
6	-1.394	-2.162	0.5228	-2.659	-5.793	-3.144	2.227	0.7522
7	-1.358	-0.1605	0.6002	-1.973	0.8823	-3.597	1.193	0.6362
8	-1.457	-0.9486	0.7394	-0.4943	-0.4982	-3.564	2.166	0.5334
9	-1.366	-1.814	0.8098	-2.344	-2.07	-3.54	1.193	0.8742
10	-1.425	0.06321	0.6188	-0.4746	-0.88	-0.04974	2.478	0.7132
11	-1.366	-3.666	0.7745	-2.039	-2.261	-2.019	1.466	0.6498
12	-1.369	-4.779	0.7515	-2.282	-3.912	-1.445	1.719	0.5189
13	-1.519	0.3	1.0	-3.0	-0.6697	-4.177	1.0	1.0
14	-1.375	-3.855	0.624	-1.615	-3.091	-1.079	2.011	0.5656
15	-1.376	-4.864	1.0	-1.381	-2.951	-1.98	1.0	1.0
16	-1.419	-4.974	1.0	-2.542	-2.482	-2.042	2.534	0.7778
17	-1.469	-4.32	0.3	-2.628	-2.757	-0.931	1.0	0.5
18	-1.378	-4.246	1.0	-1.217	-2.451	-1.794	1.763	0.5
19	-1.409	-3.904	1.0	-1.692	-1.784	-1.176	1.852	1.0
20	-1.356	-5.034	0.8484	-1.528	-2.992	-1.23	1.806	0.5
21	-1.366	-5.255	0.3	-1.323	-2.058	-1.617	1.505	0.5
22	-1.365	-2.753	0.8147	-2.132	-2.934	-3.035	1.351	0.688
23	-1.65	0.2354	0.3	-1.693	1.8	-3.051	1.0	0.5
24	-1.359	-4.182	0.9605	-2.47	-2.19	-5.291	1.743	0.6425
25	-1.377	-4.618	0.3175	-1.709	-2.825	-1.826	1.709	0.6598
26	-1.36	-4.716	0.6659	-1.086	-2.492	-1.014	1.248	0.6906
27	-1.363	-0.8163	0.831	-2.159	0.6714	-3.988	1.437	0.6369
28	-1.365	-3.208	0.8849	-2.387	-2.277	-4.094	1.469	0.7267
29	-1.491	-2.64	0.3146	-2.72	-2.235	-3.162	1.577	0.5261
30	-1.373	-2.419	0.5783	-1.985	-4.593	-0.2901	2.045	0.5732
31	-1.4	-5.969	0.4113	-2.329	0.1311	-5.538	2.738	0.6068
32	-1.499	-5.795	0.5584	-0.178	-3.244	-4.061	1.057	0.508
33	-1.457	-3.113	0.9634	-0.422	-4.206	-1.789	1.987	0.6836
34	-1.372	-2.865	0.7614	-1.505	-2.07	-4.461	1.634	0.9723
35	-1.371	-2.907	0.856	-1.397	-1.989	-4.378	1.861	0.8259
36	-1.368	-2.447	1.0	-1.7	-2.718	-3.784	1.005	0.9573
37	-1.373	-3.591	1.0	-1.59	-2.627	-3.381	1.351	0.9302
38	-1.348	-3.588	1.0	-1.909	-2.899	-4.723	1.739	0.6063
39	-1.411	-3.949	1.0	-2.045	-2.256	-4.447	2.322	1.0
40	-1.352	-3.344	1.0	-2.252	-2.499	-5.117	1.0	0.5
41	-1.363	-2.548	1.0	-2.219	-3.058	-4.758	1.618	0.5
42	-1.353	-3.447	1.0	-2.287	-3.416	-4.166	1.0	0.5
43	-1.411	-3.575	0.403	-2.635	-3.262	-5.615	1.558	0.6814
44	-1.372	-2.979	1.0	-1.585	-3.56	-3.864	1.791	0.5
45	-1.354	-4.308	1.0	-2.306	-2.561	-4.458	1.0	0.5
46	-1.365	-3.449	1.0	-1.179	-2.868	-4.695	1.0	0.5
47	-1.381	-3.527	1.0	-2.304	-3.819	-2.835	1.0	1.0
48	-1.35	-3.97	1.0	-1.909	-1.476	-4.887	1.0	0.5
49	-1.368	-4.382	1.0	-1.569	-2.171	-5.768	1.0	0.5
50	-1.494	-4.512	1.0	-3.0	-1.336	-5.316	1.0	0.5
51	-1.389	-5.461	1.0	-1.462	-1.956	-0.6024	1.824	0.5
52	-1.399	-4.266	0.5629	-1.188	-2.059	-4.868	1.38	0.6287
53	-1.387	-2.169	1.0	-2.508	-3.768	-3.724	1.0	0.9673
54	-1.359	-3.305	1.0	-1.753	-1.914	-5.561	1.692	0.5
55	-1.38	-4.22	1.0	-1.916	-3.475	-4.89	1.0	1.0
56	-1.347	-2.784	1.0	-2.061	-1.497	-4.684	1.0	0.5

```
# Extract the optimal hyperparameters from the Bayesian Optimization object
```

```
best_params_MBBR = xgb_bo_MBBR.max['params']
```

```
# Transform the hyperparameters from log space to original space
```

```
best_params_MBBR['eta'] = 10 ** best_params_MBBR['eta']
```

```
best_params_MBBR['alpha'] = 10 ** best_params_MBBR['alpha']
```

```
best_params_MBBR['lambda'] = 10 ** best_params_MBBR['lambda']
```

```
best_params_MBBR['gamma'] = 10 ** best_params_MBBR['gamma']
```

```
best_params_MBBR['max_depth'] = int(round(2 ** best_params_MBBR['max_depth']))
```

```
# Define the remaining xgboost parameters
```

```
best_params_MBBR['objective'] = 'reg:squarederror' # or 'binary:logistic' for classification
```

```
best_params_MBBR['eval_metric'] = 'rmse' # or 'auc' for classification
```

```
best_params_MBBR['colsample_bytree'] = best_params_MBBR['col_subsample']
```

```
best_params_MBBR['subsample'] = best_params_MBBR['subsample']
```

```
del best_params_MBBR['col_subsample']
del best_params_MBBR['lambda']
```

```
best_params_MBBR
```

```
{'alpha': 1.9952623149688795,
 'eta': 0.01847655149540597,
 'gamma': 4.703474432422985e-06,
 'max_depth': 2,
 'subsample': 0.5,
 'lambda': 1.9952623149688795,
 'objective': 'reg:squarederror',
 'eval_metric': 'rmse',
 'colsample_bytree': 1.0}
```

✓ For Raw Dataset

```
# Convert the data into DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_train_MBBR_dropped)
```

```
# Define the function to be optimized
def xgb_evaluate(eta, alpha, lambda, gamma, subsample, col_subsample, max_depth):
```

```
    eta = 10**eta
    alpha = 10**alpha
    lambda = 10**lambda
    gamma = 10**gamma
    max_depth = int(round(2**max_depth))
```

```
    params = {'eval_metric': 'rmse',
              'objective': 'reg:squarederror',
              'max_depth': max_depth,
              'eta': eta,
              'gamma': gamma,
              'subsample': subsample,
              'alpha': alpha,
              'lambda': lambda,
              'colsample_bytree': col_subsample,}
```

```
    cv_result = xgb.cv(params, dtrain, num_boost_round=1000, nfold=5, early_stopping_rounds=30, seed=808)
    return -1.0 * cv_result['test-rmse-mean'].iloc[-1]
```

```
# Specify the hyperparameters to be tuned
```

```
xgb_bo_MBBR_dropped = BayesianOptimization(xgb_evaluate, {'eta': (-3, 0),
                  'alpha': (-6, 0.3),
                  'lambda': (-6, 0.3),
                  'gamma': (-6, 1.8),
                  'subsample': (0.5, 1),
                  'col_subsample': (0.3, 1),
                  'max_depth': (1, 3)},
                  random_state=808)
```

```
# Optimize the hyperparameters
```

```
xgb_bo_MBBR_dropped.maximize(n_iter=1000, init_points=10)# Convert the data into DMatrix format
```

	iter	target	alpha	col_su...	eta	gamma	lambda	max_depth	subsample
1	1	-1.592	0.04075	0.4513	-2.68	-1.662	-1.582	2.026	0.7673
2	2	-1.483	-4.514	0.7529	-1.843	-2.2	-1.339	1.596	0.5436
3	3	-1.534	-1.108	0.5069	-1.136	-4.974	-0.5216	2.693	0.8202
4	4	-1.654	-3.147	0.6275	-2.063	1.604	-0.4504	1.466	0.7294
5	5	-1.566	-2.356	0.4052	-0.3806	-0.09483	-5.01	1.643	0.6674
6	6	-1.569	-2.162	0.5228	-2.659	-5.793	-3.144	2.227	0.7522
7	7	-1.551	-0.1605	0.6002	-1.973	0.8823	-3.597	1.193	0.6362
8	8	-1.504	-0.9486	0.7394	-0.4943	-0.4982	-3.564	2.166	0.5334
9	9	-1.522	-1.814	0.8098	-2.344	-2.07	-3.54	1.193	0.8742
10	10	-1.56	0.06321	0.6188	-0.4746	-0.88	-0.04974	2.478	0.7132
11	11	-1.542	-5.294	0.6658	-1.215	-5.342	-0.3629	2.421	0.9975
12	12	-1.622	-4.018	0.5023	-2.865	-0.7772	-2.167	2.849	0.9048
13	13	-1.477	-5.736	0.7323	-1.385	-3.443	-1.267	1.381	0.8479
14	14	-1.509	-4.497	0.9763	-1.066	-3.183	-1.361	1.0	0.5
15	15	-1.531	-5.569	0.6173	-2.111	-2.615	-0.5689	1.211	0.5
16	16	-1.466	-5.204	0.7135	-1.388	-2.85	-1.831	1.981	0.9342
17	17	-1.521	-5.694	0.9343	-1.793	-3.029	-2.629	1.079	1.0
18	18	-1.594	-5.566	0.3124	-0.3799	-2.642	-1.238	2.008	1.0
19	19	-1.489	-4.999	0.9307	-1.94	-3.211	-1.569	1.718	0.7673
20	20	-1.532	-5.065	0.3	-1.622	-2.747	-1.902	1.44	0.5
21	21	-1.486	-5.674	1.0	-1.506	-3.387	-1.795	2.028	1.0

22	-1.496	-4.543	1.0	-1.514	-2.659	-1.56	2.235	1.0
23	-1.487	-5.493	1.0	-1.826	-2.421	-1.567	1.919	1.0
24	-1.5	-5.182	1.0	-1.595	-2.758	-2.514	2.436	1.0
25	-1.514	-5.436	1.0	-1.351	-4.249	-1.52	1.0	1.0
26	-1.542	-5.319	0.3351	-1.932	-3.1	-1.522	2.547	1.0
27	-1.475	-5.226	1.0	-1.3	-2.865	-1.619	1.382	1.0
28	-1.529	-3.836	1.0	-1.964	-2.295	-0.9241	1.0	1.0
29	-1.494	-4.798	1.0	-0.9794	-3.405	-2.251	1.843	1.0
30	-1.478	-4.697	1.0	-1.243	-1.978	-2.177	1.594	1.0
31	-1.502	-3.63	1.0	-0.7825	-2.029	-2.309	1.527	0.5
32	-1.487	-4.715	1.0	-0.4656	-2.014	-3.38	1.625	1.0
33	-1.503	-4.725	1.0	-0.5147	-0.9607	-2.811	1.0	1.0
34	-1.538	-3.747	1.0	0.0	-2.665	-3.611	1.0	1.0
35	-1.515	-5.402	1.0	-0.801	-1.543	-2.923	2.097	0.5
36	-1.518	-4.402	1.0	-1.567	-1.757	-3.441	1.123	1.0
37	-1.538	-5.597	1.0	0.0	-1.542	-4.242	1.227	1.0
38	-1.517	-4.125	1.0	-0.165	-1.594	-2.846	2.41	1.0
39	-1.48	0.03424	1.0	0.0	-1.2	-3.918	3.0	0.5
40	-1.54	0.04954	1.0	0.0	-1.751	-3.964	1.927	0.5
41	-1.514	0.1303	1.0	0.0	-0.1083	-3.742	3.0	0.5
42	-1.555	-0.7936	1.0	0.0	-1.263	-3.339	3.0	1.0
43	-1.489	-4.593	1.0	-1.203	-1.466	-1.485	1.387	0.5
44	-1.494	-4.851	0.935	-0.1538	-2.206	-2.432	1.018	0.9766
45	-1.523	-4.05	1.0	-0.2374	-1.385	-1.646	1.0	1.0
46	-1.472	0.3	0.8367	-0.3183	-1.037	-4.812	3.0	0.5
47	-1.479	0.3	1.0	-1.099	-1.109	-4.229	3.0	0.5
48	-1.527	0.3	0.3	-0.6027	-1.715	-4.459	3.0	0.5
49	-1.476	-0.3601	1.0	-0.815	-0.5424	-4.82	3.0	0.5
50	-1.487	0.3	1.0	-1.237	-0.6958	-5.547	3.0	0.5
51	-1.557	0.06674	0.3517	-0.3089	-0.1756	-5.308	2.71	0.926
52	-1.476	-0.5299	1.0	-1.743	-1.145	-4.928	3.0	0.5
53	-1.478	6.574e-05	1.0	-2.012	-0.3042	-4.579	3.0	0.5
54	-1.505	0.3	1.0	-2.669	-1.15	-5.25	3.0	0.5
55	-1.481	-0.988	1.0	-1.731	-0.5035	-4.081	3.0	0.5
56	-1.48	-1.051	1.0	-2.212	-0.159	-5.326	3.0	0.5

```
# Extract the optimal hyperparameters from the Bayesian Optimization object
best_params_MBBR_dropped = xgb_bo_MBBR_dropped.max['params']
```

```
# Transform the hyperparameters from log space to original space
best_params_MBBR_dropped['eta'] = 10 ** best_params_MBBR_dropped['eta']
best_params_MBBR_dropped['alpha'] = 10 ** best_params_MBBR_dropped['alpha']
best_params_MBBR_dropped['lambda'] = 10 ** best_params_MBBR_dropped['lambda']
best_params_MBBR_dropped['gamma'] = 10 ** best_params_MBBR_dropped['gamma']
best_params_MBBR_dropped['max_depth'] = int(round(2 ** best_params_MBBR_dropped['max_depth']))
```

```
# Define the remaining xgboost parameters
best_params_MBBR_dropped['objective'] = 'reg:squarederror' # or 'binary:logistic' for classification
best_params_MBBR_dropped['eval_metric'] = 'rmse' # or 'auc' for classification
best_params_MBBR_dropped['colsample_bytree'] = best_params_MBBR_dropped['col_subsample']
best_params_MBBR_dropped['subsample'] = best_params_MBBR_dropped['subsample']
```

```
del best_params_MBBR_dropped['col_subsample']
del best_params_MBBR_dropped['lambda']
```

```
best_params_MBBR_dropped
```

```
{'alpha': 1.9952623149688795,
 'eta': 1.0,
 'gamma': 3.5362952331631605e-06,
 'max_depth': 8,
 'subsample': 0.5,
 'lambda': 0.0013228849154839303,
 'objective': 'reg:squarederror',
 'eval_metric': 'rmse',
 'colsample_bytree': 1.0}
```

✓ Final Model Training and Testing

✓ Optimized XGBoost 1

- Optimal Features
- Optimal Hyperparameters
- Trained on Imputed Dataset

```
# Convert test data to DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR[list(optimal_features_MBBR)], label=y_train_MBBR)
dtest = xgb.DMatrix(X_test_MBBR[list(optimal_features_MBBR)], label=y_test_MBBR)
```

▼ Determination of optimal num_boost_round

```
evals_result_MBBR = {}
```

```
# Train the final model
```

```
final_model_MBBR = xgb.train(best_params_MBBR, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'),
                                                         evals_result=evals_result_MBBR)
```

```
→ [0]    train-rmse:1.89580    test-rmse:1.67027
   [1]    train-rmse:1.88279    test-rmse:1.66121
   [2]    train-rmse:1.86983    test-rmse:1.65117
   [3]    train-rmse:1.85503    test-rmse:1.63816
   [4]    train-rmse:1.84013    test-rmse:1.62826
   [5]    train-rmse:1.82402    test-rmse:1.61516
   [6]    train-rmse:1.81098    test-rmse:1.60398
   [7]    train-rmse:1.79862    test-rmse:1.59604
   [8]    train-rmse:1.78539    test-rmse:1.58872
   [9]    train-rmse:1.77381    test-rmse:1.57905
  [10]    train-rmse:1.76347    test-rmse:1.57110
  [11]    train-rmse:1.75198    test-rmse:1.56259
  [12]    train-rmse:1.74061    test-rmse:1.55405
  [13]    train-rmse:1.72947    test-rmse:1.54406
  [14]    train-rmse:1.71715    test-rmse:1.53413
  [15]    train-rmse:1.70808    test-rmse:1.52759
  [16]    train-rmse:1.69746    test-rmse:1.51868
  [17]    train-rmse:1.68885    test-rmse:1.51230
  [18]    train-rmse:1.67976    test-rmse:1.50661
  [19]    train-rmse:1.66936    test-rmse:1.49765
  [20]    train-rmse:1.66027    test-rmse:1.49073
  [21]    train-rmse:1.65113    test-rmse:1.48462
  [22]    train-rmse:1.64232    test-rmse:1.47666
  [23]    train-rmse:1.63335    test-rmse:1.47031
  [24]    train-rmse:1.62509    test-rmse:1.46416
  [25]    train-rmse:1.61514    test-rmse:1.45730
  [26]    train-rmse:1.60729    test-rmse:1.45055
  [27]    train-rmse:1.59934    test-rmse:1.44506
  [28]    train-rmse:1.59199    test-rmse:1.44014
  [29]    train-rmse:1.58413    test-rmse:1.43469
  [30]    train-rmse:1.57770    test-rmse:1.43374
  [31]    train-rmse:1.57084    test-rmse:1.43176
  [32]    train-rmse:1.56459    test-rmse:1.43103
  [33]    train-rmse:1.55869    test-rmse:1.42501
  [34]    train-rmse:1.55168    test-rmse:1.42217
  [35]    train-rmse:1.54636    test-rmse:1.41497
  [36]    train-rmse:1.53919    test-rmse:1.41274
  [37]    train-rmse:1.53261    test-rmse:1.40841
  [38]    train-rmse:1.52630    test-rmse:1.40358
  [39]    train-rmse:1.52123    test-rmse:1.39981
  [40]    train-rmse:1.51534    test-rmse:1.39807
  [41]    train-rmse:1.51026    test-rmse:1.39440
  [42]    train-rmse:1.50429    test-rmse:1.39118
  [43]    train-rmse:1.49733    test-rmse:1.38788
  [44]    train-rmse:1.49122    test-rmse:1.38612
  [45]    train-rmse:1.48625    test-rmse:1.38300
  [46]    train-rmse:1.48070    test-rmse:1.38198
  [47]    train-rmse:1.47634    test-rmse:1.38154
  [48]    train-rmse:1.47079    test-rmse:1.37774
  [49]    train-rmse:1.46649    test-rmse:1.37562
  [50]    train-rmse:1.46387    test-rmse:1.37056
  [51]    train-rmse:1.45850    test-rmse:1.36800
  [52]    train-rmse:1.45355    test-rmse:1.36455
  [53]    train-rmse:1.45090    test-rmse:1.36019
  [54]    train-rmse:1.44614    test-rmse:1.35726
  [55]    train-rmse:1.44204    test-rmse:1.35660
  [56]    train-rmse:1.43899    test-rmse:1.35240
  [57]    train-rmse:1.43568    test-rmse:1.34738
  ...
```

```
# Train the final model
```

```
final_model_MBBR = xgb.train(best_params_MBBR, dtrain, num_boost_round=(np.argmax(evals_result_MBBR['train']['rmse'])+1), early_
                             evals_result=evals_result_MBBR)
```

```
# Make predictions on the test set
```

```
y_pred_final_MBBR = final_model_MBBR.predict(dtest)
```

```
→ [0]    train-rmse:1.89580    test-rmse:1.67027
   [1]    train-rmse:1.88279    test-rmse:1.66121
```

[2]	train-rmse:1.86983	test-rmse:1.65117
[3]	train-rmse:1.85503	test-rmse:1.63816
[4]	train-rmse:1.84013	test-rmse:1.62826
[5]	train-rmse:1.82402	test-rmse:1.61516
[6]	train-rmse:1.81098	test-rmse:1.60398
[7]	train-rmse:1.79862	test-rmse:1.59604
[8]	train-rmse:1.78539	test-rmse:1.58872
[9]	train-rmse:1.77381	test-rmse:1.57905
[10]	train-rmse:1.76347	test-rmse:1.57110
[11]	train-rmse:1.75198	test-rmse:1.56259
[12]	train-rmse:1.74061	test-rmse:1.55405
[13]	train-rmse:1.72947	test-rmse:1.54406
[14]	train-rmse:1.71715	test-rmse:1.53413
[15]	train-rmse:1.70808	test-rmse:1.52759
[16]	train-rmse:1.69746	test-rmse:1.51868
[17]	train-rmse:1.68885	test-rmse:1.51230
[18]	train-rmse:1.67976	test-rmse:1.50661
[19]	train-rmse:1.66936	test-rmse:1.49765
[20]	train-rmse:1.66027	test-rmse:1.49073
[21]	train-rmse:1.65113	test-rmse:1.48462
[22]	train-rmse:1.64232	test-rmse:1.47666
[23]	train-rmse:1.63335	test-rmse:1.47031
[24]	train-rmse:1.62509	test-rmse:1.46416
[25]	train-rmse:1.61514	test-rmse:1.45730
[26]	train-rmse:1.60729	test-rmse:1.45055
[27]	train-rmse:1.59934	test-rmse:1.44506
[28]	train-rmse:1.59199	test-rmse:1.44014
[29]	train-rmse:1.58413	test-rmse:1.43469
[30]	train-rmse:1.57770	test-rmse:1.43374
[31]	train-rmse:1.57084	test-rmse:1.43176
[32]	train-rmse:1.56459	test-rmse:1.43103
[33]	train-rmse:1.55869	test-rmse:1.42501
[34]	train-rmse:1.55168	test-rmse:1.42217
[35]	train-rmse:1.54636	test-rmse:1.41497
[36]	train-rmse:1.53919	test-rmse:1.41274
[37]	train-rmse:1.53261	test-rmse:1.40841
[38]	train-rmse:1.52630	test-rmse:1.40358
[39]	train-rmse:1.52123	test-rmse:1.39981
[40]	train-rmse:1.51534	test-rmse:1.39807
[41]	train-rmse:1.51026	test-rmse:1.39440
[42]	train-rmse:1.50429	test-rmse:1.39118
[43]	train-rmse:1.49733	test-rmse:1.38788
[44]	train-rmse:1.49122	test-rmse:1.38612
[45]	train-rmse:1.48625	test-rmse:1.38300
[46]	train-rmse:1.48070	test-rmse:1.38198
[47]	train-rmse:1.47634	test-rmse:1.38154
[48]	train-rmse:1.47079	test-rmse:1.37774
[49]	train-rmse:1.46649	test-rmse:1.37562
[50]	train-rmse:1.46387	test-rmse:1.37056
[51]	train-rmse:1.45850	test-rmse:1.36800
[52]	train-rmse:1.45355	test-rmse:1.36455
[53]	train-rmse:1.45090	test-rmse:1.36019
[54]	train-rmse:1.44614	test-rmse:1.35726
[55]	train-rmse:1.44204	test-rmse:1.35660
[56]	train-rmse:1.43899	test-rmse:1.35240
[57]	train-rmse:1.43568	test-rmse:1.34738

✓ Optimized XGBoost 2

- Optimal Features
- Optimal Hyperparameters
- Trained on Raw Dataset

```
# Convert test data to DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_train_MBBR_dropped)
dtest = xgb.DMatrix(X_test_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_test_MBBR_dropped)
```

✓ Determination of optimal num_boost_round

```
evals_result_MBBR_dropped = {}
```

```
# Train the final model
final_model_MBBR_dropped = xgb.train(best_params_MBBR_dropped, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(c
    evals_result=evals_result_MBBR_dropped)
```

```
↩ [0] train-rmse:1.40870 test-rmse:1.71140
   [1] train-rmse:1.24142 test-rmse:1.80005
   [2] train-rmse:1.23788 test-rmse:1.89925
```

[3]	train-rmse:1.17688	test-rmse:2.07116
[4]	train-rmse:1.16131	test-rmse:2.06647
[5]	train-rmse:1.14904	test-rmse:2.07194
[6]	train-rmse:1.13173	test-rmse:2.09764
[7]	train-rmse:1.09817	test-rmse:2.06489
[8]	train-rmse:1.09566	test-rmse:2.07525
[9]	train-rmse:1.09732	test-rmse:2.00322
[10]	train-rmse:1.11771	test-rmse:2.03887
[11]	train-rmse:1.11762	test-rmse:1.97840
[12]	train-rmse:1.09616	test-rmse:2.02052
[13]	train-rmse:1.10368	test-rmse:2.05023
[14]	train-rmse:1.08516	test-rmse:2.11001
[15]	train-rmse:1.08605	test-rmse:2.01925
[16]	train-rmse:1.08371	test-rmse:2.07020
[17]	train-rmse:1.10454	test-rmse:2.13831
[18]	train-rmse:1.07585	test-rmse:2.05609
[19]	train-rmse:1.08941	test-rmse:2.02826
[20]	train-rmse:1.07219	test-rmse:2.02971
[21]	train-rmse:1.08773	test-rmse:2.15171
[22]	train-rmse:1.07759	test-rmse:2.15857
[23]	train-rmse:1.07577	test-rmse:2.05925
[24]	train-rmse:1.05174	test-rmse:1.99394
[25]	train-rmse:1.04971	test-rmse:2.05482
[26]	train-rmse:1.07189	test-rmse:1.94906
[27]	train-rmse:1.05623	test-rmse:2.04827
[28]	train-rmse:1.06337	test-rmse:2.03752
[29]	train-rmse:1.06458	test-rmse:2.04442
[30]	train-rmse:1.05708	test-rmse:1.97872

```
# Train the final model
```

```
final_model_MBBR_dropped = xgb.train(best_params_MBBR_dropped, dtrain, num_boost_round=(np.argmin(
    evals_result_MBBR_dropped['train-rmse'])))
```

```
# Make predictions on the test set
```

```
y_pred_final_MBBR_dropped = final_model_MBBR_dropped.predict(dtest)
```

[0]	train-rmse:1.40870	test-rmse:1.71140
[1]	train-rmse:1.24142	test-rmse:1.80005
[2]	train-rmse:1.23788	test-rmse:1.89925
[3]	train-rmse:1.17688	test-rmse:2.07116
[4]	train-rmse:1.16131	test-rmse:2.06647
[5]	train-rmse:1.14904	test-rmse:2.07194
[6]	train-rmse:1.13173	test-rmse:2.09764
[7]	train-rmse:1.09817	test-rmse:2.06489
[8]	train-rmse:1.09566	test-rmse:2.07525
[9]	train-rmse:1.09732	test-rmse:2.00322
[10]	train-rmse:1.11771	test-rmse:2.03887
[11]	train-rmse:1.11762	test-rmse:1.97840
[12]	train-rmse:1.09616	test-rmse:2.02052
[13]	train-rmse:1.10368	test-rmse:2.05023
[14]	train-rmse:1.08516	test-rmse:2.11001
[15]	train-rmse:1.08605	test-rmse:2.01925
[16]	train-rmse:1.08371	test-rmse:2.07020
[17]	train-rmse:1.10454	test-rmse:2.13831
[18]	train-rmse:1.07585	test-rmse:2.05609
[19]	train-rmse:1.08941	test-rmse:2.02826
[20]	train-rmse:1.07219	test-rmse:2.02971
[21]	train-rmse:1.08773	test-rmse:2.15171
[22]	train-rmse:1.07759	test-rmse:2.15857
[23]	train-rmse:1.07577	test-rmse:2.05925
[24]	train-rmse:1.05174	test-rmse:1.99394
[25]	train-rmse:1.04971	test-rmse:2.05482

✧ Untuned XGBoost 1

- No Feature Selection
- No Hyperparameter Tuning
- Trained on **Imputed Dataset**

```
dtrain = xgb.DMatrix(X_train_MBBR, label=y_train_MBBR)
dtest = xgb.DMatrix(X_test_MBBR, label=y_test_MBBR)
```

```
params = {
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'seed': 808
}
```

```
# Train the out of the box xgboost model
oob_model_imputed_MBBR = xgb.train(params, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'),(dte

# Make predictions on the test set
y_pred_oob_imputed_MBBR = oob_model_imputed_MBBR.predict(dtest)
```

```
[0]    train-rmse:1.47201    test-rmse:1.51096
[1]    train-rmse:1.15157    test-rmse:1.38235
[2]    train-rmse:0.94782    test-rmse:1.30895
[3]    train-rmse:0.76046    test-rmse:1.30758
[4]    train-rmse:0.62928    test-rmse:1.28191
[5]    train-rmse:0.54018    test-rmse:1.29572
[6]    train-rmse:0.47154    test-rmse:1.29169
[7]    train-rmse:0.39595    test-rmse:1.30802
[8]    train-rmse:0.32587    test-rmse:1.29979
[9]    train-rmse:0.28197    test-rmse:1.29710
[10]   train-rmse:0.25631    test-rmse:1.30237
[11]   train-rmse:0.22741    test-rmse:1.30027
[12]   train-rmse:0.20132    test-rmse:1.29721
[13]   train-rmse:0.18149    test-rmse:1.30206
[14]   train-rmse:0.17027    test-rmse:1.29957
[15]   train-rmse:0.16005    test-rmse:1.30077
[16]   train-rmse:0.14791    test-rmse:1.30029
[17]   train-rmse:0.14193    test-rmse:1.30079
[18]   train-rmse:0.13104    test-rmse:1.30447
[19]   train-rmse:0.12352    test-rmse:1.30207
[20]   train-rmse:0.10744    test-rmse:1.30586
[21]   train-rmse:0.10091    test-rmse:1.30774
[22]   train-rmse:0.09209    test-rmse:1.30716
[23]   train-rmse:0.08606    test-rmse:1.30459
[24]   train-rmse:0.07575    test-rmse:1.30595
[25]   train-rmse:0.07116    test-rmse:1.30555
[26]   train-rmse:0.06702    test-rmse:1.30691
[27]   train-rmse:0.06049    test-rmse:1.30692
[28]   train-rmse:0.05737    test-rmse:1.30844
[29]   train-rmse:0.05287    test-rmse:1.30800
[30]   train-rmse:0.04900    test-rmse:1.30782
[31]   train-rmse:0.04693    test-rmse:1.30807
[32]   train-rmse:0.04091    test-rmse:1.30711
[33]   train-rmse:0.03893    test-rmse:1.30770
[34]   train-rmse:0.03709    test-rmse:1.30662
```

✓ Untuned XGBoost 2

- No Feature Selection
- No Hyperparameter Tuning
- Trained on **Non-Imputed (Raw) Dataset**

```
dtrain = xgb.DMatrix(X_train_MBBR_dropped, label=y_train_MBBR_dropped)
dtest = xgb.DMatrix(X_test_MBBR_dropped, label=y_test_MBBR_dropped)
```

```
params = {
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'seed': 808
}
```

```
# Train the out of the box xgboost model
oob_model_MBBR = xgb.train(params, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'),(dtest, 'tes
```

```
# Make predictions on the test set
y_pred_oob_MBBR = oob_model_MBBR.predict(dtest)
```

```
[0]    train-rmse:1.54958    test-rmse:1.48793
[1]    train-rmse:1.24651    test-rmse:1.46705
[2]    train-rmse:0.97843    test-rmse:1.47737
[3]    train-rmse:0.79490    test-rmse:1.49702
[4]    train-rmse:0.66246    test-rmse:1.51259
[5]    train-rmse:0.56802    test-rmse:1.53194
[6]    train-rmse:0.50260    test-rmse:1.53623
[7]    train-rmse:0.42722    test-rmse:1.55193
[8]    train-rmse:0.38135    test-rmse:1.56244
[9]    train-rmse:0.33632    test-rmse:1.57958
[10]   train-rmse:0.28438    test-rmse:1.58056
[11]   train-rmse:0.26198    test-rmse:1.58829
[12]   train-rmse:0.23728    test-rmse:1.60317
[13]   train-rmse:0.21035    test-rmse:1.59406
```

[14]	train-rmse:0.18668	test-rmse:1.60318
[15]	train-rmse:0.16417	test-rmse:1.59872
[16]	train-rmse:0.15322	test-rmse:1.60157
[17]	train-rmse:0.13933	test-rmse:1.60271
[18]	train-rmse:0.12366	test-rmse:1.60606
[19]	train-rmse:0.11291	test-rmse:1.60355
[20]	train-rmse:0.10587	test-rmse:1.60410
[21]	train-rmse:0.10023	test-rmse:1.60809
[22]	train-rmse:0.09571	test-rmse:1.60888
[23]	train-rmse:0.08612	test-rmse:1.61020
[24]	train-rmse:0.08365	test-rmse:1.61056
[25]	train-rmse:0.08099	test-rmse:1.61055
[26]	train-rmse:0.06964	test-rmse:1.60636
[27]	train-rmse:0.06130	test-rmse:1.60454
[28]	train-rmse:0.05457	test-rmse:1.60606
[29]	train-rmse:0.05111	test-rmse:1.60840
[30]	train-rmse:0.04773	test-rmse:1.60865
[31]	train-rmse:0.04429	test-rmse:1.60893

Naive Model 1

- **Always predicts** the mean effluent chlorine residual of the **imputed training dataset**

```
y_pred_naive_MBBR = np.full(y_test_MBBR.shape, y_train_MBBR.mean())
```

Naive Model 2

- **Always predicts** the mean effluent chlorine residual of the **Non-imputed (raw) training dataset**

```
y_pred_naive_orig_MBBR = np.full(y_test_MBBR.shape, y_train_orig_MBBR.mean())
```

Model Evaluation

```
def compute_metrics(y_pred,y_test):
    std_obs = np.std(y_test)
    std_sim = np.std(y_pred)

    mean_obs = np.mean(y_test)
    mean_sim = np.mean(y_pred)

    # Computing correlation
    r = np.corrcoef(y_test, y_pred)[0, 1]

    # Computing KGE
    alpha = std_sim / std_obs
    beta = mean_sim / mean_obs

    kge = 1 - np.sqrt(np.square(r - 1) + np.square(alpha - 1) + np.square(beta - 1))

    # PBIAS Calculation
    pbias = np.sum((y_test - y_pred)) / np.sum(y_test) * 100

    # Computing NSE
    nse = 1 - (np.sum((y_test-y_pred)**2))/(np.sum((y_test-np.mean(y_test))**2))

    if nse > 0.35:
        nse = (nse,'good')
    else:
        nse = (nse,'bad')
    if abs(pbias) < 15:
        pbias = (abs(pbias),'good')
    else:
        pbias = (abs(pbias),'bad')
    if kge > -0.41:
        kge = (kge,'good')
    else:
        kge = (kge,'bad')

    return(nse,pbias,kge)
```



```
def compute_nrmse(y_true, y_pred):
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    nrmse = rmse / (np.max(y_true) - np.min(y_true))
    return nrmse
```

```
non_imputed_mask_MBBR = ~np.isnan(y_test_orig_MBBR)
```

✓ Model Metrics evaluated on Imputed Test Set

✓ Optimized XGBoost 1

```
nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR, y_test_MBBR)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")
```

```
rmse = mean_squared_error(y_test_MBBR, y_pred_final_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_final_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Final model metrics:

```
NSE: (0.4673669183754957, 'good'),
PBIAS: (9.522618153764945, 'good'),
KGE: (0.6053104518179091, 'good')
```

```
Root Mean Squared Error: 1.227489850806859
Normalized Root Mean Squared Error: 0.24030733179460825
```

✓ Untuned XGBoost 1

```
nse_naive, pbias_naive, kge_naive = compute_metrics(y_pred_oob_imputed_MBBR, y_test_MBBR)
print(f"Final model metrics:\n\nNSE: {nse_naive}, \nPBIAS: {pbias_naive}, \nKGE: {kge_naive}")
```

```
rmse = mean_squared_error(y_test_MBBR, y_pred_oob_imputed_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_oob_imputed_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Final model metrics:

```
NSE: (0.39648104797579675, 'good'),
PBIAS: (10.577718752388224, 'good'),
KGE: (0.62675033458363, 'good')
```

```
Root Mean Squared Error: 1.306619989012038
Normalized Root Mean Squared Error: 0.25579874491230187
```

✓ Naive Model 1

```
rmse = mean_squared_error(y_test_MBBR, y_pred_naive_MBBR, squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_naive_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Root Mean Squared Error: 1.6828674290904115
Normalized Root Mean Squared Error: 0.3294572100803468

✓ Naive Model 2

```
rmse = mean_squared_error(y_test_MBBR, y_pred_naive_orig_MBBR, squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_naive_orig_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

```

↗ Root Mean Squared Error: 1.6828999269810896
Normalized Root Mean Squared Error: 0.329463572235922

```

✓ Model Metrics evaluated on Non-Imputed (Raw) Test Set

✓ Optimized XGBoost 1

```

nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR[non_imputed_mask_MBBR], y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")

```

```

rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_final_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")

```

```

nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_final_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")

```

```

↗ Final model metrics:

```

```

NSE: (0.1979687665925126, 'bad'),
PBIAS: (18.554287736175443, 'bad'),
KGE: (0.42710098316921574, 'good')

```

```

Root Mean Squared Error: 1.4822555796583785
Normalized Root Mean Squared Error: 0.29704520634436443

```

✓ Optimized XGBoost 2

```

nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR_dropped, y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")

```

```

rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_final_MBBR_dropped, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")

```

```

nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_final_MBBR_dropped)
print(f"Normalized Root Mean Squared Error: {nrmse}")

```

```

↗ Final model metrics:

```

```

NSE: (-0.5413266020400924, 'bad'),
PBIAS: (20.85801659458364, 'bad'),
KGE: (0.29327000235137435, 'good')

```

```

Root Mean Squared Error: 2.054824760065243
Normalized Root Mean Squared Error: 0.411788529071191

```

✓ Untuned XGBoost 2

```

nse_naive, pbias_naive, kge_naive = compute_metrics(y_pred_oob_MBBR, y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_naive}, \nPBIAS: {pbias_naive}, \nKGE: {kge_naive}")

```

```

rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_oob_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")

```

```

nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_oob_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")

```

```

↗ Final model metrics:

```

```

NSE: (0.05503191654014583, 'bad'),
PBIAS: (8.674598570215617, 'good'),
KGE: (0.44942410077765016, 'good')

```

```

Root Mean Squared Error: 1.608925827702348
Normalized Root Mean Squared Error: 0.3224300255916529

```

✓ Naive Model 1

```

rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_naive_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"Root Mean Squared Error: {rmse}")

```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_naive_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

```
↗ Root Mean Squared Error: 1.6631881676095377
Normalized Root Mean Squared Error: 0.33330424200591935
```

✓ Naive Model 2

```
rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_naive_orig_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_naive_orig_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

```
↗ Root Mean Squared Error: 1.6630941882756551
Normalized Root Mean Squared Error: 0.33328540847207516
```

✓ Feature Importance

```
# Get feature importance
importance_MBBR = final_model_MBBR.get_score(importance_type='gain')
```

```
name_dict_MBBR = {
    'Flow.Rate.Influent..m3.d.': 'Flow Rate Influent',
    'BOD.Influent..ppm.': 'BOD Influent',
    'Total.Coliiform.Effluent..MPN.100mL.': 'Total Coliiform Effluent',
    'pH.Pre.chlorination': 'pH Pre-Chlorination',
    'Chlorine.dosage..L.d.': 'Chlorine Dosage',
    'TSS.Pre.chlorination..ppm.': 'TSS Pre-Chlorination',
    'Total.Coliiform.Influent..MPN.100mL.': 'Total Coliiform Influent',
    'Fecal.Coliiform.Influent..MPN.100mL.': 'Fecal Coliiform Influent',
    'BOD.Pre.chlorination..ppm.': 'BOD Pre-Chlorination',
    'Fecal.Coliiform.Effluent..MPN.100mL.': 'Fecal Coliiform Effluent',
    'COD.Influent..ppm.': 'COD Influent',
    'COD.Pre.chlorination..ppm.': 'COD Pre-Chlorination',
}
```

```
# For visualization, it is better to convert it to a DataFrame
```

```
importance_df_MBBR = pd.DataFrame({
    'Feature': list(importance_MBBR.keys()),
    'Importance': list(importance_MBBR.values())
})
```


```
importance_df_MBBR['Feature'] = importance_df_MBBR['Feature'].replace(name_dict_MBBR)
```

```
# Sort the DataFrame by importance
```

```
importance_df_MBBR = importance_df_MBBR.sort_values(by='Importance', ascending=False)
```

```
# Plot feature importance
```


```
plt.figure(figsize=(10, 8))
plt.barh(importance_df_MBBR['Feature'], importance_df_MBBR['Importance'], color='skyblue')
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.title("Feature Importance")
plt.gca().invert_yaxis() # To show the highest importance at the top
plt.show()
```



✓ Data Visualization for Model Evaluation

✓ Optimized XGBoost on Imputed Test Dataset

```
# with imputation
plt.scatter(y_test_MBBR,y_pred_final_MBBR);
plt.xlabel('True Value');
plt.ylabel('Predicted Value');
```




```
# Create an x-axis range based on the length of the series/array
x = range(1, len(y_test_MBBR) + 1)
```

```
# Plotting
plt.figure(figsize=(10, 5))
plt.plot(x, y_test_MBBR, label='True', marker='o')
plt.plot(x, y_pred_final_MBBR, label='Predicted', marker='x')


# Adding labels and title
plt.xlabel('Test Datapoint')
plt.ylabel('Effluent Chlorine Residual')
plt.title('Comparison of True and Predicted Values on Imputed Test Data')
plt.legend()

# Show plot
plt.show()
```



✓ Optimized XGBoost on Non-Imputed (Raw) Test Dataset

```
# without imputation
plt.scatter(y_test_orig_MBBR[non_imputed_mask_MBBR], y_pred_final_MBBR[non_imputed_mask_MBBR])
plt.xlabel('True Value');
plt.ylabel('Predicted Value');
```




```
# Create an x-axis range based on the length of the series/array
x = range(1, len(y_test_orig_MBBR[non_imputed_mask_MBBR]) + 1)
```

```
# Plotting
plt.figure(figsize=(10, 5))
plt.plot(x, y_test_orig_MBBR[non_imputed_mask_MBBR], label='True', marker='o')
plt.plot(x, y_pred_final_MBBR[non_imputed_mask_MBBR], label='Predicted', marker='x')

# Adding labels and title
plt.xlabel('Test Datapoint')
plt.ylabel('Effluent Chlorine Residual')
plt.title('Comparison of True and Predicted Values on Raw Test Data')
plt.legend()

# Show plot
plt.show()
```



✓ Exporting Results

```
# Determine the maximum length of the columns
max_length = max(len(y_test_MBBR), len(y_test_MBBR_dropped), len(y_pred_final_MBBR), len(y_pred_final_MBBR_dropped), len(y_pred_

# Function to extend a series or array to the maximum length with NaN values
def extend_with_nan(data, length):
    if isinstance(data, np.ndarray):
        data = pd.Series(data)
```