

```
!pip install bayesian-optimization
```

```
Collecting bayesian-optimization
  Downloading bayesian_optimization-1.4.3-py3-none-any.whl (18 kB)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.25.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.11.4)
Requirement already satisfied: scikit-learn>=0.18.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization)
Collecting colorama>=0.4.6 (from bayesian-optimization)
  Downloading colorama-0.4.6-py2.py3-none-any.whl (25 kB)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (3.2.0)
Installing collected packages: colorama, bayesian-optimization
Successfully installed bayesian-optimization-1.4.3 colorama-0.4.6
```

```
!git clone https://github.com/808ss/thesis.git
```

```
Cloning into 'thesis'...
remote: Enumerating objects: 27, done.
remote: Counting objects: 100% (27/27), done.
remote: Compressing objects: 100% (26/26), done.
remote: Total 27 (delta 0), reused 0 (delta 0), pack-reused 0
Receiving objects: 100% (27/27), 311.32 KiB | 1.80 MiB/s, done.
```

```
import numpy as np
import pandas as pd
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import seaborn as sns
import itertools
from bayes_opt import BayesianOptimization
```

```
random_seed = 808
np.random.seed(random_seed)
```

## ✓ MBBR

### ✓ Importing MBBR and Splitting

```
MBBR = pd.read_csv('thesis/MBBR-Chlorination.csv')
MBBR.drop(columns='Date', inplace=True)
```

```
X_orig_MBBR = MBBR.drop(columns='Residual chlorine\n(ppm)')
y_orig_MBBR = MBBR['Residual chlorine\n(ppm)']
X_train_orig_MBBR, X_test_orig_MBBR, y_train_orig_MBBR, y_test_orig_MBBR = train_test_split(X_orig_MBBR,
                                                                                          y_orig_MBBR,
                                                                                          test_size = 0.3,
                                                                                          random_state=808)
```

```
df_train_orig_MBBR = pd.concat([X_train_orig_MBBR, y_train_orig_MBBR], axis=1)
df_test_orig_MBBR = pd.concat([X_test_orig_MBBR, y_test_orig_MBBR], axis=1)
```

### ✓ Data Analysis for Raw Dataset

```
missing_rate_MBBR = [(MBBR.isnull().sum()[val]/MBBR.shape[0])*100 for val in range(0, MBBR.shape[1])]
```

```
pd.options.display.float_format = '{:,.2f}'.format
MBBR_transposed = MBBR.describe().T
MBBR_transposed['Missingness Rate'] = missing_rate_MBBR
```

```
MBBR_transposed
```



|                                     | count  | mean           | std            | min        | 25%           | 50%           | 75%            | max              | Missingness Rate |
|-------------------------------------|--------|----------------|----------------|------------|---------------|---------------|----------------|------------------|------------------|
| Flow Rate Influent (m3/d)           | 332.00 | 4,787.53       | 2,211.48       | 197.00     | 3,344.00      | 4,709.50      | 6,232.00       | 11,147.00        | 0.0%             |
| Total Coliform Influent (MPN/100mL) | 270.00 | 290,896,939.26 | 733,941,441.61 | 1,600.00   | 17,250,000.00 | 40,500,000.00 | 160,000,000.00 | 5,200,000,000.00 | 18.6%            |
| Total Coliform Effluent (MPN/100mL) | 329.00 | 733,375.06     | 9,402,984.69   | 0.00       | 2.00          | 10.00         | 471.00         | 143,900,000.00   | 0.9%             |
| Fecal Coliform Influent (MPN/100mL) | 103.00 | 236,377,087.38 | 621,705,589.64 | 230,000.00 | 8,550,000.00  | 23,000,000.00 | 37,650,000.00  | 3,000,000,000.00 | 68.9%            |
| Fecal Coliform Effluent (MPN/100mL) | 171.00 | 746.87         | 3,947.85       | 2.00       | 10.00         | 10.00         | 10.00          | 24,196.00        | 48.4%            |
| BOD Influent (ppm)                  | 273.00 | 152.40         | 148.02         | 8.00       | 68.00         | 119.00        | 196.00         | 1,425.00         | 17.7%            |
| BOD Pre-chlorination\n(ppm)         | 274.00 | 11.28          | 12.82          | 1.00       | 4.00          | 8.00          | 14.00          | 119.00           | 17.4%            |

▼ Data Analysis for Training Set (Pre-Imputation)

```
missing_rate_train_orig_MBBR = [(df_train_orig_MBBR.isnull().sum()[val]/df_train_orig_MBBR.shape[0])*100 for val in range(0,df_train_orig_MBBR.shape[0])]
```

```
pd.options.display.float_format = '{:,.2f}'.format
#pd.set_option('display.float_format', '{:e}'.format)
df_train_orig_MBBR_transposed = df_train_orig_MBBR.describe().T
df_train_orig_MBBR_transposed['Missingness Rate'] = missing_rate_train_orig_MBBR
```

```
df_train_orig_MBBR_transposed
```



|                                     | count  | mean           | std            | min        | 25%           | 50%           | 75%            | max              | Missingness Rate |
|-------------------------------------|--------|----------------|----------------|------------|---------------|---------------|----------------|------------------|------------------|
| Flow Rate Influent (m3/d)           | 232.00 | 4,882.74       | 2,204.80       | 197.00     | 3,344.00      | 4,762.00      | 6,349.00       | 10,999.00        | 0.0%             |
| Total Coliform Influent (MPN/100mL) | 185.00 | 315,522,010.81 | 790,769,090.66 | 16,000.00  | 18,000,000.00 | 41,000,000.00 | 160,000,000.00 | 5,200,000,000.00 | 20.2%            |
| Total Coliform Effluent (MPN/100mL) | 230.00 | 1,045,242.88   | 11,238,969.87  | 0.00       | 2.25          | 10.00         | 1,280.75       | 143,900,000.00   | 0.8%             |
| Fecal Coliform Influent (MPN/100mL) | 68.00  | 298,124,117.65 | 669,480,560.27 | 230,000.00 | 10,400,000.00 | 24,000,000.00 | 40,950,000.00  | 2,600,000,000.00 | 70.6%            |
| Fecal Coliform Effluent (MPN/100mL) | 120.00 | 892.02         | 4,352.93       | 2.00       | 10.00         | 10.00         | 10.00          | 24,196.00        | 48.2%            |
| BOD Influent (ppm)                  | 187.00 | 162.30         | 167.60         | 8.00       | 70.50         | 122.00        | 199.00         | 1,425.00         | 19.4%            |
| BOD Pre-chlorination\n(ppm)         | 188.00 | 11.12          | 12.31          | 1.00       | 5.00          | 8.00          | 14.00          | 119.00           | 18.9%            |

▼ Data Analysis for Testing Set (Pre-imputation)

```
missing_rate_test_orig_MBBR = [(df_test_orig_MBBR.isnull().sum()[val]/df_test_orig_MBBR.shape[0])*100 for val in range(0,df_test_orig_MBBR.shape[0])]
```

```
#pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
df_test_orig_MBBR_transposed = df_test_orig_MBBR.describe().T
df_test_orig_MBBR_transposed['Missingness Rate'] = missing_rate_test_orig_MBBR
```

```
df_test_orig_MBBR_transposed
```



|                                     | count        | mean         | std          | min          | 25%          | 50%          | 75%          | max          | Missing  |
|-------------------------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|----------|
| Flow Rate Influent (m3/d)           | 1.000000e+02 | 4.566650e+03 | 2.222244e+03 | 2.170000e+02 | 3.343750e+03 | 4.641500e+03 | 6.112250e+03 | 1.114700e+04 | 0.000000 |
| Total Coliform Influent (MPN/100mL) | 8.500000e+01 | 2.373012e+08 | 5.924907e+08 | 1.600000e+03 | 1.700000e+07 | 4.000000e+07 | 1.600000e+08 | 3.500000e+09 | 1.500000 |
| Total Coliform Effluent (MPN/100mL) | 9.900000e+01 | 8.833667e+03 | 3.764360e+04 | 0.000000e+00 | 2.000000e+00 | 1.000000e+01 | 9.000000e+01 | 2.419600e+05 | 1.000000 |
| Fecal Coliform Influent (MPN/100mL) | 3.500000e+01 | 1.164114e+08 | 5.038721e+08 | 1.000000e+06 | 7.450000e+06 | 1.700000e+07 | 3.045000e+07 | 3.000000e+09 | 6.500000 |
| Fecal Coliform Effluent (MPN/100mL) | 5.100000e+01 | 4.053529e+02 | 2.779390e+03 | 2.000000e+00 | 2.000000e+00 | 1.000000e+01 | 1.000000e+01 | 1.986300e+04 | 4.900000 |
| BOD Influent (ppm)                  | 8.600000e+01 | 1.308605e+02 | 8.921358e+01 | 1.900000e+01 | 6.600000e+01 | 1.060000e+02 | 1.830000e+02 | 4.090000e+02 | 1.400000 |
| BOD Pre-chlorination\n(ppm)         | 8.600000e+01 | 1.162791e+01 | 1.393434e+01 | 1.000000e+00 | 4.000000e+00 | 8.000000e+00 | 1.400000e+01 | 1.080000e+02 | 1.400000 |

## ▼ Data Imputation

### ▼ Exporting Datasets to R

```
df_train_orig_MBBR.to_csv('MBBR_train_set.csv',index=False)
df_test_orig_MBBR.to_csv('MBBR_test_set.csv',index=False)
```

```
# Export to R for mixgb
```

### ▼ Mixgb imputation

```
1 library(mixgb)
2 library(openxlsx)
3 set.seed(808)
4
5 MBBR_train_set <- read.csv("C:/Users/nikko/PycharmProjects/Thesis/MBBR_train_set.csv")
6 MBBR_test_set <- read.csv("C:/Users/nikko/PycharmProjects/Thesis/MBBR_test_set.csv")
7
8 MBBR_train_set_df = as.data.frame(MBBR_train_set)
9 MBBR_test_set_df = as.data.frame(MBBR_test_set)
10
11 clean_MBBR_train_set_df <- data_clean(MBBR_train_set_df)
12 clean_MBBR_test_set_df <- data_clean(MBBR_test_set_df)
13
14 cv.results_1 <- mixgb_cv(data = clean_MBBR_train_set_df, nrounds = 5000, verbose = FALSE)
15 cv.results_1$evaluation.log
16 cv.results_1$best.nrounds
17
18 mixgb_obj <- mixgb(data = clean_MBBR_train_set_df, m = 5, nrounds = cv.results_1$best.nrounds, save.models = TRUE)
19 MBBR_train_imputed <- mixgb_obj$imputed.data
20
21 MBBR_test_imputed <- impute_new(object = mixgb_obj, newdata = clean_MBBR_test_set_df)
22
23 write.xlsx(MBBR_train_imputed[[1]], file = 'mbbr_m1_imputed_train.xlsx')
24 write.xlsx(MBBR_train_imputed[[2]], file = 'mbbr_m2_imputed_train.xlsx')
25 write.xlsx(MBBR_train_imputed[[3]], file = 'mbbr_m3_imputed_train.xlsx')
26 write.xlsx(MBBR_train_imputed[[4]], file = 'mbbr_m4_imputed_train.xlsx')
27 write.xlsx(MBBR_train_imputed[[5]], file = 'mbbr_m5_imputed_train.xlsx')
28
29 write.xlsx(MBBR_test_imputed[[1]], file = 'mbbr_m1_imputed_test.xlsx')
30 write.xlsx(MBBR_test_imputed[[2]], file = 'mbbr_m2_imputed_test.xlsx')
31 write.xlsx(MBBR_test_imputed[[3]], file = 'mbbr_m3_imputed_test.xlsx')
32 write.xlsx(MBBR_test_imputed[[4]], file = 'mbbr_m4_imputed_test.xlsx')
33 write.xlsx(MBBR_test_imputed[[5]], file = 'mbbr_m5_imputed_test.xlsx')
```

### ▼ Import imputed datasets from R

```
dfs = []
for val in range(1,6):
    source = f'thesis/mbbr_m{val}_imputed_train.xlsx'
```

```
dfs.append(pd.read_excel(source))

average_MBBR_train = pd.concat(dfs).groupby(level=0).mean()

dfs = []
for val in range(1,6):
    source = f'thesis/mbbr_m{val}_imputed_test.xlsx'
    dfs.append(pd.read_excel(source))

average_MBBR_test = pd.concat(dfs).groupby(level=0).mean()
```

▼ Data Analysis for Training Set (Post-Imputation)

```
#pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
average_MBBR_train_transposed = average_MBBR_train.describe().T
```

average\_MBBR\_train\_transposed

|                                     | count        | mean         | std          | min          | 25%          | 50%          | 75%          |         |
|-------------------------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|---------|
| Flow.Rate.Influent..m3.d.           | 2.320000e+02 | 4.882741e+03 | 2.204801e+03 | 1.970000e+02 | 3.344000e+03 | 4.762000e+03 | 6.349000e+03 | 1.09990 |
| Total.Coliform.Influent..MPN.100mL. | 2.320000e+02 | 3.480306e+08 | 7.784700e+08 | 1.600000e+04 | 2.175000e+07 | 5.400000e+07 | 2.200000e+08 | 5.20000 |
| Total.Coliform.Effluent..MPN.100mL. | 2.320000e+02 | 1.036336e+06 | 1.119062e+07 | 0.000000e+00 | 2.750000e+00 | 1.000000e+01 | 1.600000e+03 | 1.43900 |
| Fecal.Coliform.Influent..MPN.100mL. | 2.320000e+02 | 1.965840e+08 | 5.001635e+08 | 2.300000e+05 | 1.428500e+07 | 2.740000e+07 | 3.821000e+07 | 2.60000 |
| Fecal.Coliform.Effluent..MPN.100mL. | 2.320000e+02 | 4.182178e+03 | 8.256544e+03 | 2.000000e+00 | 8.800000e+00 | 1.000000e+01 | 2.282500e+02 | 2.41960 |
| BOD.Influent..ppm.                  | 2.320000e+02 | 1.585655e+02 | 1.533443e+02 | 8.000000e+00 | 7.585000e+01 | 1.250000e+02 | 1.950500e+02 | 1.42500 |
| BOD.Pre.chlorination..ppm.          | 2.320000e+02 | 1.199138e+01 | 1.246310e+01 | 1.000000e+00 | 5.000000e+00 | 9.000000e+00 | 1.500000e+01 | 1.19000 |
| COD.Influent..ppm.                  | 2.320000e+02 | 3.458931e+02 | 5.439009e+02 | 1.300000e+01 | 1.750000e+02 | 2.510000e+02 | 3.745000e+02 | 7.73400 |
| COD.Pre.chlorination..ppm.          | 2.320000e+02 | 4.940431e+01 | 3.795535e+01 | 5.000000e+00 | 2.475000e+01 | 4.150000e+01 | 6.300000e+01 | 3.43000 |
| TSS.Pre.chlorination..ppm.          | 2.320000e+02 | 1.756897e+01 | 2.094736e+01 | 1.000000e+00 | 6.000000e+00 | 1.200000e+01 | 2.000000e+01 | 1.60000 |
| pH.Pre.chlorination                 | 2.320000e+02 | 7.185121e+00 | 3.010609e-01 | 6.120000e+00 | 7.000000e+00 | 7.200000e+00 | 7.362500e+00 | 8.38000 |
| Chlorine.dosage..L.d.               | 2.320000e+02 | 8.721017e+02 | 4.852711e+02 | 0.000000e+00 | 6.000000e+02 | 8.500000e+02 | 1.160000e+03 | 2.80000 |
| Residual.chlorine..ppm.             | 2.320000e+02 | 2.082141e+00 | 1.913163e+00 | 0.000000e+00 | 3.815000e-01 | 1.205700e+00 | 3.917150e+00 | 5.48000 |

▼ Data Analysis for Testing Set (Post-Imputation)

```
pd.options.display.float_format = '{:,.2f}'.format
pd.set_option('display.float_format', '{:e}'.format)
average_MBBR_test_transposed = average_MBBR_test.describe().T
```

average\_MBBR\_test\_transposed



|                                      | count  | mean           | std            | min          | 25%           | 50%           | 75%            |               |
|--------------------------------------|--------|----------------|----------------|--------------|---------------|---------------|----------------|---------------|
| Flow.Rate.Influent..m3.d.            | 100.00 | 4,566.65       | 2,222.24       | 217.00       | 3,343.75      | 4,641.50      | 6,112.25       | 11,147        |
| Total.Coliiform.Influent..MPN.100mL. | 100.00 | 288,973,956.00 | 647,069,268.14 | 1,600.00     | 22,750,000.00 | 45,000,000.00 | 200,000,000.00 | 3,500,000,000 |
| Total.Coliiform.Effluent..MPN.100mL. | 100.00 | 8,933.07       | 37,466.19      | 0.00         | 2.00          | 10.00         | 134.75         | 241,960       |
| Fecal.Coliiform.Influent..MPN.100mL. | 100.00 | 147,284,740.00 | 427,084,482.81 | 1,000,000.00 | 12,000,000.00 | 25,320,000.00 | 41,560,000.00  | 3,000,000,000 |
| Fecal.Coliiform.Effluent..MPN.100mL. | 100.00 | 2,363.05       | 6,219.15       | 2.00         | 10.00         | 10.00         | 107.30         | 24,196        |
| BOD.Influent..ppm.                   | 100.00 | 134.13         | 88.52          | 19.00        | 66.45         | 108.40        | 188.25         | 408           |
| BOD.Pre.chlorination..ppm.           | 100.00 | 12.06          | 13.39          | 1.00         | 5.00          | 9.00          | 15.00          | 108           |
| COD.Influent..ppm.                   | 100.00 | 268.66         | 192.32         | 41.00        | 135.75        | 210.00        | 357.25         | 1,256         |
| COD.Pre.chlorination..ppm.           | 100.00 | 45.03          | 47.09          | 5.00         | 20.00         | 29.70         | 53.55          | 314           |
| TSS.Pre.chlorination..ppm.           | 100.00 | 18.00          | 24.39          | 1.00         | 5.00          | 10.00         | 19.25          | 164           |
| pH.Pre.chlorination                  | 100.00 | 7.21           | 0.37           | 5.28         | 7.10          | 7.21          | 7.40           | 8             |
| Chlorine.dosage..L.d.                | 100.00 | 809.32         | 498.65         | 0.00         | 488.70        | 748.90        | 1,005.05       | 2,900         |
| Residual.chlorine..ppm.              | 100.00 | 2.14           | 1.69           | 0.01         | 0.49          | 2.05          | 3.57           | 5             |

## Exhaustive Feature Selection

### For Imputed Dataset

```
pd.reset_option('display.float_format')
```

```
X_train_MBBR = average_MBBR_train.drop(columns=['Residual.chlorine..ppm.', 'Total.Coliiform.Effluent..MPN.100mL.', 'Fecal.Coliiform.
y_train_MBBR = average_MBBR_train['Residual.chlorine..ppm.']
X_test_MBBR = average_MBBR_test.drop(columns=['Residual.chlorine..ppm.', 'Total.Coliiform.Effluent..MPN.100mL.', 'Fecal.Coliiform.Ef
y_test_MBBR = average_MBBR_test['Residual.chlorine..ppm.'])
```

```
features_wo_chlorine_dosage = X_train_MBBR.columns[:-1]
features_wo_chlorine_dosage
```



```
Index(['Flow.Rate.Influent..m3.d.', 'Total.Coliiform.Influent..MPN.100mL.',
      'Fecal.Coliiform.Influent..MPN.100mL.', 'BOD.Influent..ppm.',
      'BOD.Pre.chlorination..ppm.', 'COD.Influent..ppm.',
      'COD.Pre.chlorination..ppm.', 'TSS.Pre.chlorination..ppm.',
      'pH.Pre.chlorination'],
      dtype='object')
```

```
# Generate all combinations of the other features
combinations = []
for r in range(1, len(features_wo_chlorine_dosage) + 1):
    combinations.extend(itertools.combinations(features_wo_chlorine_dosage, r))
```

```
# Add the first feature to each combination
combinations = [(X_train_MBBR.columns[-1],) + combo for combo in combinations]
```

```
params = {'objective': 'reg:squarederror'}
```

```
results = []
for combo in combinations:
    dtrain = xgb.DMatrix(X_train_MBBR[list(combo)], label=y_train_MBBR)
    cv_result = xgb.cv(params, dtrain, num_boost_round=10, nfold=5, metrics='rmse', seed=808)
    last_round_metrics = cv_result.iloc[-1]
    results.append([combo, last_round_metrics['train-rmse-mean'], last_round_metrics['test-rmse-mean'],
                  last_round_metrics['train-rmse-std'], last_round_metrics['test-rmse-std']])
```

```
results_df_MBBR = pd.DataFrame(results, columns=['Combination', 'Train RMSE', 'Validation RMSE', 'Train RMSE Std. Dev.', 'Valid
```

```
results_df_MBBR.sort_values(by='Validation RMSE')
```



|     | Combination  | Train RMSE | Validation RMSE | Train RMSE Std. Dev. | Validation RMSE Std. Dev |
|-----|--|------------|-----------------|----------------------|--------------------------|
| 474 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.386445   | 1.611471        | 0.028575             | 0.143105                 |
| 350 | (Chlorine.dosage..L.d., Total.Coliiform.Influen... | 0.410715   | 1.622629        | 0.044093             | 0.123522                 |
| 409 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.362284   | 1.631725        | 0.047728             | 0.139558                 |
| 410 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.382969   | 1.633061        | 0.036756             | 0.105161                 |
| 278 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.374995   | 1.635504        | 0.032118             | 0.089523                 |
| ... | ...  | ...        | ...             | ...                  | ...                      |
| 106 | (Chlorine.dosage..L.d., Fecal.Coliiform.Influen... | 0.788887   | 2.085864        | 0.043836             | 0.179699                 |
| 100 | (Chlorine.dosage..L.d., Fecal.Coliiform.Influen... | 0.742502   | 2.099164        | 0.053870             | 0.108392                 |
| 2   | (Chlorine.dosage..L.d., Fecal.Coliiform.Influen... | 0.966307   | 2.111154        | 0.030007             | 0.200713                 |
| 6   | (Chlorine.dosage..L.d., COD.Pre.chlorination.....  | 1.070077   | 2.117816        | 0.101896             | 0.117162                 |
| 27  | (Chlorine.dosage..L.d., Fecal.Coliiform.Influen... | 0.857085   | 2.160103        | 0.053470             | 0.113003                 |

511 rows x 5 columns

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[0:3]
```



|     | Combination  | Train RMSE | Validation RMSE | Train RMSE Std. Dev. | Validation RMSE Std. Dev |
|-----|--|------------|-----------------|----------------------|--------------------------|
| 474 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.386445   | 1.611471        | 0.028575             | 0.143105                 |
| 350 | (Chlorine.dosage..L.d., Total.Coliiform.Influen... | 0.410715   | 1.622629        | 0.044093             | 0.123522                 |
| 409 | (Chlorine.dosage..L.d., Flow.Rate.Influent..m3...  | 0.362284   | 1.631725        | 0.047728             | 0.139558                 |

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[0]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'Fecal.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'COD.Pre.chlorination..ppm.',
 'TSS.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[1]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'BOD.Pre.chlorination..ppm.',
 'TSS.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
results_df_MBBR.sort_values(by='Validation RMSE').iloc[2]['Combination']
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'COD.Influent..ppm.',
 'TSS.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
optimal_features_MBBR = results_df_MBBR.sort_values(by='Validation RMSE').iloc[0]['Combination']
optimal_features_MBBR
```



```
('Chlorine.dosage..L.d.',
 'Flow.Rate.Influent..m3.d.',
 'Total.Coliiform.Influent..MPN.100mL.',
 'Fecal.Coliiform.Influent..MPN.100mL.',
 'BOD.Influent..ppm.',
 'COD.Pre.chlorination..ppm.',
 'TSS.Pre.chlorination..ppm.',
 'pH.Pre.chlorination')
```

```
results_df_MBBR['count'] = results_df_MBBR['Combination'].apply(lambda x: len(x))
results_df_MBBR.to_csv('MBBR Exhaustive Feature Selection.csv', index=False)
```

## ▼ For Raw Dataset

```
non_imputed_mask_MBBR_train = ~np.isnan(y_train_orig_MBBR)
non_imputed_mask_MBBR_test = ~np.isnan(y_test_orig_MBBR)
```

```
X_train_MBBR_dropped = X_train_orig_MBBR[non_imputed_mask_MBBR_train]
y_train_MBBR_dropped = y_train_orig_MBBR[non_imputed_mask_MBBR_train]
X_test_MBBR_dropped = X_test_orig_MBBR[non_imputed_mask_MBBR_test]
y_test_MBBR_dropped = y_test_orig_MBBR[non_imputed_mask_MBBR_test]
```

```
features_wo_chlorine_dosage_dropped = X_train_MBBR_dropped.columns[:-1]
features_wo_chlorine_dosage_dropped
```

```
Index(['Flow Rate Influent (m3/d)', 'Total Coliform Influent (MPN/100mL)',
      'Total Coliform Effluent (MPN/100mL)',
      'Fecal Coliform Influent (MPN/100mL)',
      'Fecal Coliform Effluent (MPN/100mL)', 'BOD Influent (ppm)',
      'BOD Pre-chlorination\n(ppm)', 'COD Influent (ppm)',
      'COD Pre-chlorination\n(ppm)', 'TSS Pre-chlorination (ppm)',
      'pH Pre-chlorination'],
      dtype='object')
```

```
# Generate all combinations of the other features
combinations = []
for r in range(1, len(features_wo_chlorine_dosage_dropped) + 1):
    combinations.extend(itertools.combinations(features_wo_chlorine_dosage_dropped, r))
```

```
# Add the first feature to each combination
combinations = [(X_train_MBBR_dropped.columns[-1],) + combo for combo in combinations]
```

```
params = {'objective': 'reg:squarederror'}
```

```
results = []
for combo in combinations:
    dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(combo)], label=y_train_MBBR_dropped)
    cv_result = xgb.cv(params, dtrain, num_boost_round=10, nfold=5, metrics='rmse', seed=808)
    last_round_metrics = cv_result.iloc[-1]
    results.append([combo, last_round_metrics['train-rmse-mean'], last_round_metrics['test-rmse-mean'],
                  last_round_metrics['train-rmse-std'], last_round_metrics['test-rmse-std']])
```

```
results_df_MBBR_dropped = pd.DataFrame(results, columns=['Combination', 'Train RMSE', 'Validation RMSE', 'Train RMSE Std. Dev.',
```

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE')
```

```
↗
```

|      | Combination                                       | Train RMSE | Validation RMSE | Train RMSE Std. Dev. | Validation RMSE Std. Dev |
|------|---|------------|-----------------|----------------------|--------------------------|
| 31   | (Chlorine dosage (L/d), Total Coliform Effluen... | 1.022659   | 1.512094        | 0.031956             | 0.180507                 |
| 2    | (Chlorine dosage (L/d), Total Coliform Effluen... | 1.031525   | 1.520116        | 0.038504             | 0.184722                 |
| 32   | (Chlorine dosage (L/d), Total Coliform Effluen... | 0.636437   | 1.554748        | 0.050116             | 0.192632                 |
| 160  | (Chlorine dosage (L/d), Total Coliform Effluen... | 0.612772   | 1.574227        | 0.047039             | 0.187918                 |
| 574  | (Chlorine dosage (L/d), Flow Rate Influent (m3... | 0.424320   | 1.577279        | 0.044534             | 0.284987                 |
| ...  | ...   | ...        | ...             | ...                  | ...                      |
| 1000 | (Chlorine dosage (L/d), Fecal Coliform Influen... | 0.420868   | 2.110952        | 0.038620             | 0.188357                 |
| 139  | (Chlorine dosage (L/d), Total Coliform Influen... | 0.576095   | 2.111329        | 0.039815             | 0.206058                 |
| 44   | (Chlorine dosage (L/d), Fecal Coliform Influen... | 0.804355   | 2.112326        | 0.033380             | 0.167510                 |
| 525  | (Chlorine dosage (L/d), Fecal Coliform Influen... | 0.415596   | 2.112704        | 0.030191             | 0.144628                 |
| 58   | (Chlorine dosage (L/d), BOD Pre-chlorination\n... | 0.648844   | 2.153476        | 0.061506             | 0.111363                 |

2047 rows x 5 columns

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0:3]
```

|    | Combination                                       | Train RMSE | Validation RMSE | Train RMSE Std. Dev. | Validation RMSE Std. Dev |
|----|---|------------|-----------------|----------------------|--------------------------|
| 31 | (Chlorine dosage (L/d), Total Coliform Effluen... | 1.022659   | 1.512094        | 0.031956             | 0.180507                 |
| 2  | (Chlorine dosage (L/d), Total Coliform Effluen... | 1.031525   | 1.520116        | 0.038504             | 0.184722                 |
| 32 | (Chlorine dosage (L/d), Total Coliform Effluen... | 0.636437   | 1.554748        | 0.050116             | 0.192632                 |

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0]['Combination']
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'Fecal Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[1]['Combination']
```

```
('Chlorine dosage (L/d)', 'Total Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[2]['Combination']
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'BOD Influent (ppm)')
```

```
optimal_features_MBBR_dropped = results_df_MBBR_dropped.sort_values(by='Validation RMSE').iloc[0]['Combination']
optimal_features_MBBR_dropped
```

```
('Chlorine dosage (L/d)',
 'Total Coliform Effluent (MPN/100mL)',
 'Fecal Coliform Effluent (MPN/100mL)')
```

```
results_df_MBBR_dropped['count'] = results_df_MBBR_dropped['Combination'].apply(lambda x: len(x))
results_df_MBBR_dropped.to_csv('MBBR Dropped Exhaustive Feature Selection.csv', index=False)
```

## ✓ Hyperparameter Optimization

### ✓ For Imputed Dataset

```
# Convert the data into DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR[list(optimal_features_MBBR)], label=y_train_MBBR)

# Define the function to be optimized
def xgb_evaluate(eta, alpha, lambd, gamma, subsample, col_subsample, max_depth):
    eta = 10**eta
    alpha = 10**alpha
    lambd = 10**lambd
    gamma = 10**gamma
    max_depth = int(round(2**max_depth))

    params = {'eval_metric': 'rmse',
              'objective': 'reg:squarederror',
              'max_depth': max_depth,
              'eta': eta,
              'gamma': gamma,
              'subsample': subsample,
              'alpha': alpha,
              'lambda': lambd,
              'colsample_bytree': col_subsample,}

    cv_result = xgb.cv(params, dtrain, num_boost_round=1000, nfold=5, early_stopping_rounds=30, seed=808)
    return -1.0 * cv_result['test-rmse-mean'].iloc[-1]

# Specify the hyperparameters to be tuned
xgb_bo_MBBR = BayesianOptimization(xgb_evaluate, {'eta': (-3, 0),
                                                  'alpha': (-6, 0.3),
                                                  'lambd': (-6, 0.3),
                                                  'gamma': (-6, 1.8),
                                                  'subsample': (0.5, 1),
                                                  'col_subsample': (0.3, 1),
                                                  'max_depth': (1, 3)},
```



random\_state=808)

# Optimize the hyperparameters

xgb\_bo\_MBBR.maximize(n\_iter=1000, init\_points=10)# Convert the data into DMatrix format

| iter | target | alpha    | col_su... | eta     | gamma    | lambda   | max_depth | subsample |
|------|--------|----------|-----------|---------|----------|----------|-----------|-----------|
| 1    | -1.671 | 0.04075  | 0.4513    | -2.68   | -1.662   | -1.582   | 2.026     | 0.7673    |
| 2    | -1.623 | -4.514   | 0.7529    | -1.843  | -2.2     | -1.339   | 1.596     | 0.5436    |
| 3    | -1.698 | -1.108   | 0.5069    | -1.136  | -4.974   | -0.5216  | 2.693     | 0.8202    |
| 4    | -1.771 | -3.147   | 0.6275    | -2.063  | 1.604    | -0.4504  | 1.466     | 0.7294    |
| 5    | -1.713 | -2.356   | 0.4052    | -0.3806 | -0.09483 | -5.01    | 1.643     | 0.6674    |
| 6    | -1.641 | -2.162   | 0.5228    | -2.659  | -5.793   | -3.144   | 2.227     | 0.7522    |
| 7    | -1.639 | -0.1605  | 0.6002    | -1.973  | 0.8823   | -3.597   | 1.193     | 0.6362    |
| 8    | -1.753 | -0.9486  | 0.7394    | -0.4943 | -0.4982  | -3.564   | 2.166     | 0.5334    |
| 9    | -1.63  | -1.814   | 0.8098    | -2.344  | -2.07    | -3.54    | 1.193     | 0.8742    |
| 10   | -1.711 | 0.06321  | 0.6188    | -0.4746 | -0.88    | -0.04974 | 2.478     | 0.7132    |
| 11   | -1.603 | -4.543   | 0.8022    | -1.469  | -2.3     | -1.735   | 1.546     | 0.6257    |
| 12   | -1.627 | -3.644   | 0.8309    | -1.947  | -2.894   | -2.698   | 1.318     | 0.7972    |
| 13   | -1.574 | -5.741   | 0.9132    | -1.013  | -3.386   | -2.238   | 1.468     | 0.6447    |
| 14   | -2.136 | -6.0     | 0.3       | 0.0     | -2.883   | -2.26    | 3.0       | 1.0       |
| 15   | -1.629 | -5.073   | 1.0       | -1.774  | -3.319   | -2.118   | 1.0       | 0.5056    |
| 16   | -1.646 | -6.0     | 1.0       | -1.161  | -4.403   | -2.505   | 1.0       | 0.5       |
| 17   | -1.642 | -6.0     | 1.0       | -1.186  | -2.958   | -1.212   | 1.0       | 0.5       |
| 18   | -1.64  | -5.877   | 1.0       | -1.4    | -2.537   | -3.009   | 1.0       | 0.5       |
| 19   | -1.736 | -4.575   | 1.0       | -0.3881 | -3.405   | -2.28    | 1.0       | 0.5       |
| 20   | -1.632 | -6.0     | 1.0       | -2.145  | -3.377   | -2.195   | 1.728     | 1.0       |
| 21   | -1.648 | -4.42    | 1.0       | -2.332  | -1.536   | -2.603   | 1.222     | 1.0       |
| 22   | -1.657 | -2.758   | 0.3       | -2.073  | -1.874   | -1.948   | 1.33      | 0.5024    |
| 23   | -1.678 | -2.295   | 0.6771    | -2.886  | -3.659   | -3.156   | 1.774     | 0.8583    |
| 24   | -1.702 | -4.588   | 0.3       | -2.654  | -2.756   | -2.516   | 2.168     | 0.5       |
| 25   | -1.647 | -3.348   | 1.0       | -2.132  | -2.168   | -4.241   | 1.0       | 1.0       |
| 26   | -1.651 | -3.92    | 1.0       | -1.015  | -1.276   | -1.424   | 1.0       | 1.0       |
| 27   | -1.739 | -2.486   | 1.0       | -3.0    | -0.6419  | -3.416   | 1.0       | 1.0       |
| 28   | -1.643 | -2.206   | 1.0       | -1.118  | -2.9     | -3.281   | 1.0       | 1.0       |
| 29   | -1.62  | -1.667   | 1.0       | -1.982  | -3.038   | -4.85    | 1.0       | 0.5       |
| 30   | -1.663 | -0.5516  | 0.3       | -1.985  | -3.064   | -3.99    | 1.0       | 1.0       |
| 31   | -1.601 | -2.362   | 1.0       | -1.616  | -2.875   | -4.533   | 2.355     | 0.5       |
| 32   | -1.629 | -2.764   | 1.0       | -1.445  | -3.912   | -5.073   | 1.665     | 1.0       |
| 33   | -1.625 | -2.082   | 1.0       | -2.575  | -2.819   | -5.585   | 2.3       | 1.0       |
| 34   | -1.612 | -1.253   | 1.0       | -1.557  | -3.846   | -5.092   | 2.635     | 0.5       |
| 35   | -1.728 | -1.893   | 0.3       | -1.0    | -2.943   | -5.815   | 2.107     | 0.5       |
| 36   | -1.622 | -2.381   | 1.0       | -2.442  | -4.07    | -4.731   | 3.0       | 0.5       |
| 37   | -1.645 | -1.209   | 1.0       | -2.381  | -2.9     | -4.242   | 3.0       | 0.9127    |
| 38   | -1.679 | -1.939   | 1.0       | -1.058  | -4.332   | -3.825   | 2.582     | 0.5       |
| 39   | -1.649 | -1.425   | 1.0       | -2.842  | -4.374   | -5.314   | 1.773     | 0.5       |
| 40   | -1.625 | -3.656   | 0.9717    | -2.879  | -2.655   | -5.139   | 2.935     | 0.8396    |
| 41   | -1.682 | -3.277   | 1.0       | -2.978  | -3.507   | -5.151   | 1.508     | 0.5       |
| 42   | -1.716 | -1.811   | 1.0       | -1.807  | -5.112   | -5.898   | 3.0       | 1.0       |
| 43   | -1.72  | -3.027   | 1.0       | -2.088  | -2.407   | -3.931   | 3.0       | 1.0       |
| 44   | -1.613 | -0.5171  | 1.0       | -2.554  | -2.935   | -5.305   | 2.073     | 0.5       |
| 45   | -1.622 | 0.2591   | 1.0       | -2.118  | -3.942   | -4.831   | 2.747     | 0.5       |
| 46   | -1.78  | -0.8012  | 0.3       | -3.0    | -3.605   | -5.333   | 3.0       | 0.5       |
| 47   | -1.605 | -0.9969  | 1.0       | -1.601  | -2.966   | -4.455   | 2.017     | 0.5       |
| 48   | -1.609 | -0.04149 | 1.0       | -1.554  | -3.734   | -5.27    | 1.507     | 0.5       |
| 49   | -1.621 | 0.09739  | 0.8789    | -1.223  | -2.633   | -5.104   | 2.702     | 0.9513    |
| 50   | -1.618 | 0.3      | 1.0       | -1.98   | -2.207   | -5.376   | 1.293     | 0.5       |
| 51   | -1.652 | 0.3      | 1.0       | -0.5831 | -3.943   | -4.469   | 2.491     | 0.5       |
| 52   | -1.734 | -0.9158  | 1.0       | -2.969  | -2.0     | -5.482   | 1.0       | 1.0       |
| 53   | -1.605 | 0.3      | 1.0       | -2.124  | -2.78    | -4.309   | 2.066     | 0.5       |
| 54   | -1.686 | 0.3      | 1.0       | -2.8    | -3.998   | -4.605   | 1.153     | 0.5       |
| 55   | -1.627 | 0.3      | 1.0       | -2.317  | -1.671   | -5.077   | 2.77      | 0.5       |
| 56   | -1.634 | -4.956   | 0.8517    | -2.793  | -3.61    | -5.846   | 2.85      | 0.753     |

```
# Extract the optimal hyperparameters from the Bayesian Optimization object
best_params_MBBR = xgb_bo_MBBR.max['params']

# Transform the hyperparameters from log space to original space
best_params_MBBR['eta'] = 10 ** best_params_MBBR['eta']
best_params_MBBR['alpha'] = 10 ** best_params_MBBR['alpha']
best_params_MBBR['lambda'] = 10 ** best_params_MBBR['lambda']
best_params_MBBR['gamma'] = 10 ** best_params_MBBR['gamma']
best_params_MBBR['max_depth'] = int(round(2 ** best_params_MBBR['max_depth']))

# Define the remaining xgboost parameters
best_params_MBBR['objective'] = 'reg:squarederror' # or 'binary:logistic' for classification
best_params_MBBR['eval_metric'] = 'rmse' # or 'auc' for classification
best_params_MBBR['colsample_bytree'] = best_params_MBBR['col_subsample']
best_params_MBBR['subsample'] = best_params_MBBR['subsample']

del best_params_MBBR['col_subsample']
del best_params_MBBR['lambda']
```

```
best_params_MBBR
```

```
{'alpha': 1.8143466498648387e-06,
 'eta': 0.0970581980025668,
 'gamma': 0.00041160003988771333,
 'max_depth': 3,
 'subsample': 0.6447492980647876,
 'lambda': 0.005780282239906642,
 'objective': 'reg:squarederror',
 'eval_metric': 'rmse',
 'colsample_bytree': 0.9131916743594339}
```

## ▼ For Raw Dataset

```
# Convert the data into DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_train_MBBR_dropped)

# Define the function to be optimized
def xgb_evaluate(eta, alpha, lambd, gamma, subsample, col_subsample, max_depth):
    eta = 10**eta
    alpha = 10**alpha
    lambd = 10**lambd
    gamma = 10**gamma
    max_depth = int(round(2**max_depth))

    params = {'eval_metric': 'rmse',
              'objective': 'reg:squarederror',
              'max_depth': max_depth,
              'eta': eta,
              'gamma': gamma,
              'subsample': subsample,
              'alpha': alpha,
              'lambda': lambd,
              'colsample_bytree': col_subsample,}

    cv_result = xgb.cv(params, dtrain, num_boost_round=1000, nfold=5, early_stopping_rounds=30, seed=808)
    return -1.0 * cv_result['test-rmse-mean'].iloc[-1]

# Specify the hyperparameters to be tuned
xgb_bo_MBBR_dropped = BayesianOptimization(xgb_evaluate, {'eta': (-3, 0),
                                                         'alpha': (-6, 0.3),
                                                         'lambd': (-6, 0.3),
                                                         'gamma': (-6, 1.8),
                                                         'subsample': (0.5, 1),
                                                         'col_subsample': (0.3, 1),
                                                         'max_depth': (1, 3)},
                                           random_state=808)

# Optimize the hyperparameters
xgb_bo_MBBR_dropped.maximize(n_iter=1000, init_points=10)# Convert the data into DMatrix format
```

| iter | target | alpha   | col_su... | eta    | gamma  | lambd   | max_depth | subsample |
|------|--------|---------|-----------|--------|--------|---------|-----------|-----------|
| 1    | -1.592 | 0.04075 | 0.4513    | -2.68  | -1.662 | -1.582  | 2.026     | 0.7673    |
| 2    | -1.483 | -4.514  | 0.7529    | -1.843 | -2.2   | -1.339  | 1.596     | 0.5436    |
| 3    | -1.534 | -1.108  | 0.5069    | -1.136 | -4.974 | -0.5216 | 2.693     | 0.8202    |

|    |        |           |        |         |          |          |       |        |
|----|--------|-----------|--------|---------|----------|----------|-------|--------|
| 4  | -1.654 | -3.147    | 0.6275 | -2.063  | 1.604    | -0.4504  | 1.466 | 0.7294 |
| 5  | -1.566 | -2.356    | 0.4052 | -0.3806 | -0.09483 | -5.01    | 1.643 | 0.6674 |
| 6  | -1.569 | -2.162    | 0.5228 | -2.659  | -5.793   | -3.144   | 2.227 | 0.7522 |
| 7  | -1.551 | -0.1605   | 0.6002 | -1.973  | 0.8823   | -3.597   | 1.193 | 0.6362 |
| 8  | -1.504 | -0.9486   | 0.7394 | -0.4943 | -0.4982  | -3.564   | 2.166 | 0.5334 |
| 9  | -1.522 | -1.814    | 0.8098 | -2.344  | -2.07    | -3.54    | 1.193 | 0.8742 |
| 10 | -1.56  | 0.06321   | 0.6188 | -0.4746 | -0.88    | -0.04974 | 2.478 | 0.7132 |
| 11 | -1.542 | -5.294    | 0.6658 | -1.215  | -5.342   | -0.3629  | 2.421 | 0.9975 |
| 12 | -1.622 | -4.018    | 0.5023 | -2.865  | -0.7772  | -2.167   | 2.849 | 0.9048 |
| 13 | -1.477 | -5.736    | 0.7323 | -1.385  | -3.443   | -1.267   | 1.381 | 0.8479 |
| 14 | -1.509 | -4.497    | 0.9763 | -1.066  | -3.183   | -1.361   | 1.0   | 0.5    |
| 15 | -1.531 | -5.569    | 0.6173 | -2.111  | -2.615   | -0.5689  | 1.211 | 0.5    |
| 16 | -1.466 | -5.204    | 0.7135 | -1.388  | -2.85    | -1.831   | 1.981 | 0.9342 |
| 17 | -1.521 | -5.694    | 0.9343 | -1.793  | -3.029   | -2.629   | 1.079 | 1.0    |
| 18 | -1.594 | -5.566    | 0.3124 | -0.3799 | -2.642   | -1.238   | 2.008 | 1.0    |
| 19 | -1.489 | -4.999    | 0.9307 | -1.94   | -3.211   | -1.569   | 1.718 | 0.7673 |
| 20 | -1.532 | -5.065    | 0.3    | -1.622  | -2.747   | -1.902   | 1.44  | 0.5    |
| 21 | -1.486 | -5.674    | 1.0    | -1.506  | -3.387   | -1.795   | 2.028 | 1.0    |
| 22 | -1.496 | -4.543    | 1.0    | -1.514  | -2.659   | -1.56    | 2.235 | 1.0    |
| 23 | -1.487 | -5.493    | 1.0    | -1.826  | -2.421   | -1.567   | 1.919 | 1.0    |
| 24 | -1.5   | -5.182    | 1.0    | -1.595  | -2.758   | -2.514   | 2.436 | 1.0    |
| 25 | -1.514 | -5.436    | 1.0    | -1.351  | -4.249   | -1.52    | 1.0   | 1.0    |
| 26 | -1.542 | -5.319    | 0.3351 | -1.932  | -3.1     | -1.522   | 2.547 | 1.0    |
| 27 | -1.475 | -5.226    | 1.0    | -1.3    | -2.865   | -1.619   | 1.382 | 1.0    |
| 28 | -1.529 | -3.836    | 1.0    | -1.964  | -2.295   | -0.9241  | 1.0   | 1.0    |
| 29 | -1.494 | -4.798    | 1.0    | -0.9794 | -3.405   | -2.251   | 1.843 | 1.0    |
| 30 | -1.478 | -4.697    | 1.0    | -1.243  | -1.978   | -2.177   | 1.594 | 1.0    |
| 31 | -1.502 | -3.63     | 1.0    | -0.7825 | -2.029   | -2.309   | 1.527 | 0.5    |
| 32 | -1.487 | -4.715    | 1.0    | -0.4656 | -2.014   | -3.38    | 1.625 | 1.0    |
| 33 | -1.503 | -4.725    | 1.0    | -0.5147 | -0.9607  | -2.811   | 1.0   | 1.0    |
| 34 | -1.538 | -3.747    | 1.0    | 0.0     | -2.665   | -3.611   | 1.0   | 1.0    |
| 35 | -1.515 | -5.402    | 1.0    | -0.801  | -1.543   | -2.923   | 2.097 | 0.5    |
| 36 | -1.518 | -4.402    | 1.0    | -1.567  | -1.757   | -3.441   | 1.123 | 1.0    |
| 37 | -1.538 | -5.597    | 1.0    | 0.0     | -1.542   | -4.242   | 1.227 | 1.0    |
| 38 | -1.517 | -4.125    | 1.0    | -0.165  | -1.594   | -2.846   | 2.41  | 1.0    |
| 39 | -1.48  | 0.03424   | 1.0    | 0.0     | -1.2     | -3.918   | 3.0   | 0.5    |
| 40 | -1.54  | 0.04954   | 1.0    | 0.0     | -1.751   | -3.964   | 1.927 | 0.5    |
| 41 | -1.514 | 0.1303    | 1.0    | 0.0     | -0.1083  | -3.742   | 3.0   | 0.5    |
| 42 | -1.555 | -0.7936   | 1.0    | 0.0     | -1.263   | -3.339   | 3.0   | 1.0    |
| 43 | -1.489 | -4.593    | 1.0    | -1.203  | -1.466   | -1.485   | 1.387 | 0.5    |
| 44 | -1.494 | -4.851    | 0.935  | -0.1538 | -2.206   | -2.432   | 1.018 | 0.9766 |
| 45 | -1.523 | -4.05     | 1.0    | -0.2374 | -1.385   | -1.646   | 1.0   | 1.0    |
| 46 | -1.472 | 0.3       | 0.8367 | -0.3183 | -1.037   | -4.812   | 3.0   | 0.5    |
| 47 | -1.479 | 0.3       | 1.0    | -1.099  | -1.109   | -4.229   | 3.0   | 0.5    |
| 48 | -1.527 | 0.3       | 0.3    | -0.6027 | -1.715   | -4.459   | 3.0   | 0.5    |
| 49 | -1.476 | -0.3601   | 1.0    | -0.815  | -0.5424  | -4.82    | 3.0   | 0.5    |
| 50 | -1.487 | 0.3       | 1.0    | -1.237  | -0.6958  | -5.547   | 3.0   | 0.5    |
| 51 | -1.557 | 0.06674   | 0.3517 | -0.3089 | -0.1756  | -5.308   | 2.71  | 0.926  |
| 52 | -1.476 | -0.5299   | 1.0    | -1.743  | -1.145   | -4.928   | 3.0   | 0.5    |
| 53 | -1.478 | 6.574e-05 | 1.0    | -2.012  | -0.3042  | -4.579   | 3.0   | 0.5    |
| 54 | -1.505 | 0.3       | 1.0    | -2.669  | -1.15    | -5.25    | 3.0   | 0.5    |
| 55 | -1.481 | -0.988    | 1.0    | -1.731  | -0.5035  | -4.081   | 3.0   | 0.5    |

```
# Extract the optimal hyperparameters from the Bayesian Optimization object
best_params_MBBR_dropped = xgb_bo_MBBR_dropped.max['params']
```

```
# Transform the hyperparameters from log space to original space
best_params_MBBR_dropped['eta'] = 10 ** best_params_MBBR_dropped['eta']
best_params_MBBR_dropped['alpha'] = 10 ** best_params_MBBR_dropped['alpha']
best_params_MBBR_dropped['lambda'] = 10 ** best_params_MBBR_dropped['lambda']
best_params_MBBR_dropped['gamma'] = 10 ** best_params_MBBR_dropped['gamma']
best_params_MBBR_dropped['max_depth'] = int(round(2 ** best_params_MBBR_dropped['max_depth']))
```

```
# Define the remaining xgboost parameters
best_params_MBBR_dropped['objective'] = 'reg:squarederror' # or 'binary:logistic' for classification
best_params_MBBR_dropped['eval_metric'] = 'rmse' # or 'auc' for classification
best_params_MBBR_dropped['colsample_bytree'] = best_params_MBBR_dropped['col_subsample']
best_params_MBBR_dropped['subsample'] = best_params_MBBR_dropped['subsample']
```

```
del best_params_MBBR_dropped['col_subsample']
del best_params_MBBR_dropped['lambda']
```

```
best_params_MBBR_dropped
```

```
{'alpha': 1.9952623149688795,
 'eta': 1.0,
 'gamma': 3.5362952331631605e-06,
 'max_depth': 8,
 'subsample': 0.5,
 'lambda': 0.0013228849154839303,
 'objective': 'reg:squarederror',
```

```
'eval_metric': 'rmse',
'colsample_bytree': 1.0}
```

## ✓ Final Model Training and Testing

### ✓ Optimized XGBoost 1

- Optimal Features
- Optimal Hyperparameters
- Trained on Imputed Dataset

```
# Convert test data to DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR[list(optimal_features_MBBR)], label=y_train_MBBR)
dtest = xgb.DMatrix(X_test_MBBR[list(optimal_features_MBBR)], label=y_test_MBBR)
```

### ✓ Determination of optimal num\_boost\_round

```
evals_result_MBBR = {}
```

```
# Train the final model
final_model_MBBR = xgb.train(best_params_MBBR, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'),
                                                         evals_result=evals_result_MBBR])
```

```
→ [0]    train-rmse:1.83281    test-rmse:1.62397
   [1]    train-rmse:1.80099    test-rmse:1.60884
   [2]    train-rmse:1.74685    test-rmse:1.55963
   [3]    train-rmse:1.69571    test-rmse:1.54121
   [4]    train-rmse:1.63720    test-rmse:1.49432
   [5]    train-rmse:1.60046    test-rmse:1.46383
   [6]    train-rmse:1.56640    test-rmse:1.44605
   [7]    train-rmse:1.53273    test-rmse:1.42124
   [8]    train-rmse:1.49641    test-rmse:1.39748
   [9]    train-rmse:1.46776    test-rmse:1.37692
  [10]    train-rmse:1.44477    test-rmse:1.37356
  [11]    train-rmse:1.41625    test-rmse:1.36399
  [12]    train-rmse:1.39234    test-rmse:1.35255
  [13]    train-rmse:1.37290    test-rmse:1.35734
  [14]    train-rmse:1.34499    test-rmse:1.34116
  [15]    train-rmse:1.32109    test-rmse:1.32393
  [16]    train-rmse:1.29797    test-rmse:1.31431
  [17]    train-rmse:1.28107    test-rmse:1.31604
  [18]    train-rmse:1.25722    test-rmse:1.32725
  [19]    train-rmse:1.23970    test-rmse:1.33870
  [20]    train-rmse:1.23085    test-rmse:1.33786
  [21]    train-rmse:1.21392    test-rmse:1.34890
  [22]    train-rmse:1.19262    test-rmse:1.34346
  [23]    train-rmse:1.17355    test-rmse:1.33278
  [24]    train-rmse:1.15916    test-rmse:1.33643
  [25]    train-rmse:1.14377    test-rmse:1.33581
  [26]    train-rmse:1.13682    test-rmse:1.32983
  [27]    train-rmse:1.12799    test-rmse:1.32334
  [28]    train-rmse:1.12069    test-rmse:1.32276
  [29]    train-rmse:1.11284    test-rmse:1.31973
  [30]    train-rmse:1.10153    test-rmse:1.32310
  [31]    train-rmse:1.09153    test-rmse:1.32107
  [32]    train-rmse:1.08087    test-rmse:1.32216
  [33]    train-rmse:1.07602    test-rmse:1.31794
  [34]    train-rmse:1.06127    test-rmse:1.31838
  [35]    train-rmse:1.04700    test-rmse:1.32250
  [36]    train-rmse:1.02933    test-rmse:1.33290
  [37]    train-rmse:1.01959    test-rmse:1.33235
  [38]    train-rmse:1.00896    test-rmse:1.33122
  [39]    train-rmse:1.00218    test-rmse:1.33058
  [40]    train-rmse:0.99702    test-rmse:1.33031
  [41]    train-rmse:0.98924    test-rmse:1.32241
  [42]    train-rmse:0.98024    test-rmse:1.32625
  [43]    train-rmse:0.96780    test-rmse:1.33809
  [44]    train-rmse:0.95884    test-rmse:1.34195
  [45]    train-rmse:0.95332    test-rmse:1.34101
  [46]    train-rmse:0.93901    test-rmse:1.34511
```

```
# Train the final model
final_model_MBBR = xgb.train(best_params_MBBR, dtrain, num_boost_round=(np.argmin(evals_result_MBBR['train']['rmse'])+1), early_
```

```
evals_result=evals_result_MBBR)
```

```
# Make predictions on the test set
y_pred_final_MBBR = final_model_MBBR.predict(dtest)
```

|      |                    |                   |
|------|--------------------|-------------------|
| [0]  | train-rmse:1.83281 | test-rmse:1.62397 |
| [1]  | train-rmse:1.80099 | test-rmse:1.60884 |
| [2]  | train-rmse:1.74685 | test-rmse:1.55963 |
| [3]  | train-rmse:1.69571 | test-rmse:1.54121 |
| [4]  | train-rmse:1.63720 | test-rmse:1.49432 |
| [5]  | train-rmse:1.60046 | test-rmse:1.46383 |
| [6]  | train-rmse:1.56640 | test-rmse:1.44605 |
| [7]  | train-rmse:1.53273 | test-rmse:1.42124 |
| [8]  | train-rmse:1.49641 | test-rmse:1.39748 |
| [9]  | train-rmse:1.46776 | test-rmse:1.37692 |
| [10] | train-rmse:1.44477 | test-rmse:1.37356 |
| [11] | train-rmse:1.41625 | test-rmse:1.36399 |
| [12] | train-rmse:1.39234 | test-rmse:1.35255 |
| [13] | train-rmse:1.37290 | test-rmse:1.35734 |
| [14] | train-rmse:1.34499 | test-rmse:1.34116 |
| [15] | train-rmse:1.32109 | test-rmse:1.32393 |
| [16] | train-rmse:1.29797 | test-rmse:1.31431 |
| [17] | train-rmse:1.28107 | test-rmse:1.31604 |
| [18] | train-rmse:1.25722 | test-rmse:1.32725 |
| [19] | train-rmse:1.23970 | test-rmse:1.33870 |
| [20] | train-rmse:1.23085 | test-rmse:1.33786 |
| [21] | train-rmse:1.21392 | test-rmse:1.34890 |
| [22] | train-rmse:1.19262 | test-rmse:1.34346 |
| [23] | train-rmse:1.17355 | test-rmse:1.33278 |
| [24] | train-rmse:1.15916 | test-rmse:1.33643 |
| [25] | train-rmse:1.14377 | test-rmse:1.33581 |
| [26] | train-rmse:1.13682 | test-rmse:1.32983 |
| [27] | train-rmse:1.12799 | test-rmse:1.32334 |
| [28] | train-rmse:1.12069 | test-rmse:1.32276 |
| [29] | train-rmse:1.11284 | test-rmse:1.31973 |
| [30] | train-rmse:1.10153 | test-rmse:1.32310 |
| [31] | train-rmse:1.09153 | test-rmse:1.32107 |
| [32] | train-rmse:1.08087 | test-rmse:1.32216 |
| [33] | train-rmse:1.07602 | test-rmse:1.31794 |
| [34] | train-rmse:1.06127 | test-rmse:1.31838 |
| [35] | train-rmse:1.04700 | test-rmse:1.32250 |
| [36] | train-rmse:1.02933 | test-rmse:1.33290 |
| [37] | train-rmse:1.01959 | test-rmse:1.33235 |
| [38] | train-rmse:1.00896 | test-rmse:1.33122 |
| [39] | train-rmse:1.00218 | test-rmse:1.33058 |
| [40] | train-rmse:0.99702 | test-rmse:1.33031 |
| [41] | train-rmse:0.98924 | test-rmse:1.32241 |
| [42] | train-rmse:0.98024 | test-rmse:1.32625 |
| [43] | train-rmse:0.96780 | test-rmse:1.33809 |
| [44] | train-rmse:0.95884 | test-rmse:1.34195 |
| [45] | train-rmse:0.95332 | test-rmse:1.34101 |
| [46] | train-rmse:0.93901 | test-rmse:1.34511 |

## ✓ Optimized XGBoost 2

- Optimal Features
- Optimal Hyperparameters
- Trained on Raw Dataset

```
# Convert test data to DMatrix format
dtrain = xgb.DMatrix(X_train_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_train_MBBR_dropped)
dtest = xgb.DMatrix(X_test_MBBR_dropped[list(optimal_features_MBBR_dropped)], label=y_test_MBBR_dropped)
```

## ✓ Determination of optimal num\_boost\_round

```
evals_result_MBBR_dropped = {}
```

```
# Train the final model
final_model_MBBR_dropped = xgb.train(best_params_MBBR_dropped, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(c
    evals_result=evals_result_MBBR_dropped)
```

|     |                    |                   |
|-----|--------------------|-------------------|
| [0] | train-rmse:1.40870 | test-rmse:1.71140 |
| [1] | train-rmse:1.24142 | test-rmse:1.80005 |
| [2] | train-rmse:1.23788 | test-rmse:1.89925 |
| [3] | train-rmse:1.17688 | test-rmse:2.07116 |
| [4] | train-rmse:1.16131 | test-rmse:2.06647 |
| [5] | train-rmse:1.14904 | test-rmse:2.07194 |

|      |                    |                   |
|------|--------------------|-------------------|
| [6]  | train-rmse:1.13173 | test-rmse:2.09764 |
| [7]  | train-rmse:1.09817 | test-rmse:2.06489 |
| [8]  | train-rmse:1.09566 | test-rmse:2.07525 |
| [9]  | train-rmse:1.09732 | test-rmse:2.00322 |
| [10] | train-rmse:1.11771 | test-rmse:2.03887 |
| [11] | train-rmse:1.11762 | test-rmse:1.97840 |
| [12] | train-rmse:1.09616 | test-rmse:2.02052 |
| [13] | train-rmse:1.10368 | test-rmse:2.05023 |
| [14] | train-rmse:1.08516 | test-rmse:2.11001 |
| [15] | train-rmse:1.08605 | test-rmse:2.01925 |
| [16] | train-rmse:1.08371 | test-rmse:2.07020 |
| [17] | train-rmse:1.10454 | test-rmse:2.13831 |
| [18] | train-rmse:1.07585 | test-rmse:2.05609 |
| [19] | train-rmse:1.08941 | test-rmse:2.02826 |
| [20] | train-rmse:1.07219 | test-rmse:2.02971 |
| [21] | train-rmse:1.08773 | test-rmse:2.15171 |
| [22] | train-rmse:1.07759 | test-rmse:2.15857 |
| [23] | train-rmse:1.07577 | test-rmse:2.05925 |
| [24] | train-rmse:1.05174 | test-rmse:1.99394 |
| [25] | train-rmse:1.04971 | test-rmse:2.05482 |
| [26] | train-rmse:1.07189 | test-rmse:1.94906 |
| [27] | train-rmse:1.05623 | test-rmse:2.04827 |
| [28] | train-rmse:1.06337 | test-rmse:2.03752 |
| [29] | train-rmse:1.06458 | test-rmse:2.04442 |
| [30] | train-rmse:1.05708 | test-rmse:1.97872 |

```
# Train the final model
```

```
final_model_MBBR_dropped = xgb.train(best_params_MBBR_dropped, dtrain, num_boost_round=(np.argmin(evals_result_MBBR_dropped['train-rmse'])), evals_result=evals_result_MBBR_dropped)
```

```
# Make predictions on the test set
```

```
y_pred_final_MBBR_dropped = final_model_MBBR_dropped.predict(dtest)
```

|      |                    |                   |
|------|--------------------|-------------------|
| [0]  | train-rmse:1.40870 | test-rmse:1.71140 |
| [1]  | train-rmse:1.24142 | test-rmse:1.80005 |
| [2]  | train-rmse:1.23788 | test-rmse:1.89925 |
| [3]  | train-rmse:1.17688 | test-rmse:2.07116 |
| [4]  | train-rmse:1.16131 | test-rmse:2.06647 |
| [5]  | train-rmse:1.14904 | test-rmse:2.07194 |
| [6]  | train-rmse:1.13173 | test-rmse:2.09764 |
| [7]  | train-rmse:1.09817 | test-rmse:2.06489 |
| [8]  | train-rmse:1.09566 | test-rmse:2.07525 |
| [9]  | train-rmse:1.09732 | test-rmse:2.00322 |
| [10] | train-rmse:1.11771 | test-rmse:2.03887 |
| [11] | train-rmse:1.11762 | test-rmse:1.97840 |
| [12] | train-rmse:1.09616 | test-rmse:2.02052 |
| [13] | train-rmse:1.10368 | test-rmse:2.05023 |
| [14] | train-rmse:1.08516 | test-rmse:2.11001 |
| [15] | train-rmse:1.08605 | test-rmse:2.01925 |
| [16] | train-rmse:1.08371 | test-rmse:2.07020 |
| [17] | train-rmse:1.10454 | test-rmse:2.13831 |
| [18] | train-rmse:1.07585 | test-rmse:2.05609 |
| [19] | train-rmse:1.08941 | test-rmse:2.02826 |
| [20] | train-rmse:1.07219 | test-rmse:2.02971 |
| [21] | train-rmse:1.08773 | test-rmse:2.15171 |
| [22] | train-rmse:1.07759 | test-rmse:2.15857 |
| [23] | train-rmse:1.07577 | test-rmse:2.05925 |
| [24] | train-rmse:1.05174 | test-rmse:1.99394 |
| [25] | train-rmse:1.04971 | test-rmse:2.05482 |

## ✖ Untuned XGBoost 1

- No Feature Selection
- No Hyperparameter Tuning
- Trained on **Imputed Dataset**

```
dtrain = xgb.DMatrix(X_train_MBBR, label=y_train_MBBR)
dtest = xgb.DMatrix(X_test_MBBR, label=y_test_MBBR)
```

```
params = {
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'seed': 808
}
```

```
# Train the out of the box xgboost model
```

```
oob_model_imputed_MBBR = xgb.train(params, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'), (dtest, 'test')])
```

```
# Make predictions on the test set
y_pred_oob_imputed_MBBR = oob_model_imputed_MBBR.predict(dtest)
```

```
[0]    train-rmse:1.50822    test-rmse:1.51609
[1]    train-rmse:1.20441    test-rmse:1.50001
[2]    train-rmse:0.98990    test-rmse:1.49717
[3]    train-rmse:0.82211    test-rmse:1.50169
[4]    train-rmse:0.71358    test-rmse:1.52169
[5]    train-rmse:0.62480    test-rmse:1.50629
[6]    train-rmse:0.54690    test-rmse:1.49833
[7]    train-rmse:0.48596    test-rmse:1.49060
[8]    train-rmse:0.42910    test-rmse:1.50414
[9]    train-rmse:0.38121    test-rmse:1.50953
[10]   train-rmse:0.33375    test-rmse:1.50861
[11]   train-rmse:0.29921    test-rmse:1.50841
[12]   train-rmse:0.26334    test-rmse:1.51815
[13]   train-rmse:0.23309    test-rmse:1.51110
[14]   train-rmse:0.22078    test-rmse:1.50781
[15]   train-rmse:0.20472    test-rmse:1.51623
[16]   train-rmse:0.19130    test-rmse:1.51380
[17]   train-rmse:0.17399    test-rmse:1.51761
[18]   train-rmse:0.16732    test-rmse:1.52052
[19]   train-rmse:0.15921    test-rmse:1.52261
[20]   train-rmse:0.14906    test-rmse:1.52428
[21]   train-rmse:0.13491    test-rmse:1.51987
[22]   train-rmse:0.13105    test-rmse:1.52192
[23]   train-rmse:0.12136    test-rmse:1.52082
[24]   train-rmse:0.11906    test-rmse:1.52181
[25]   train-rmse:0.10825    test-rmse:1.52442
[26]   train-rmse:0.09492    test-rmse:1.52286
[27]   train-rmse:0.08642    test-rmse:1.52315
[28]   train-rmse:0.08302    test-rmse:1.52431
[29]   train-rmse:0.07634    test-rmse:1.52492
[30]   train-rmse:0.07197    test-rmse:1.52554
[31]   train-rmse:0.06772    test-rmse:1.52645
[32]   train-rmse:0.06204    test-rmse:1.52908
[33]   train-rmse:0.05586    test-rmse:1.52984
[34]   train-rmse:0.05126    test-rmse:1.53033
[35]   train-rmse:0.04565    test-rmse:1.52998
[36]   train-rmse:0.04242    test-rmse:1.52988
```

## ✓ Untuned XGBoost 2

- No Feature Selection
- No Hyperparameter Tuning
- Trained on **Non-Imputed (Raw) Dataset**

```
dtrain = xgb.DMatrix(X_train_MBBR_dropped, label=y_train_MBBR_dropped)
dtest = xgb.DMatrix(X_test_MBBR_dropped, label=y_test_MBBR_dropped)
```

```
params = {
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'seed': 808
}
```

```
# Train the out of the box xgboost model
```

```
oob_model_MBBR = xgb.train(params, dtrain, num_boost_round=1000, early_stopping_rounds=30, evals=[(dtrain, 'train'), (dtest, 'tes
```

```
# Make predictions on the test set
```

```
y_pred_oob_MBBR = oob_model_MBBR.predict(dtest)
```

```
[0]    train-rmse:1.54958    test-rmse:1.48793
[1]    train-rmse:1.24651    test-rmse:1.46705
[2]    train-rmse:0.97843    test-rmse:1.47737
[3]    train-rmse:0.79490    test-rmse:1.49702
[4]    train-rmse:0.66246    test-rmse:1.51259
[5]    train-rmse:0.56802    test-rmse:1.53194
[6]    train-rmse:0.50260    test-rmse:1.53623
[7]    train-rmse:0.42722    test-rmse:1.55193
[8]    train-rmse:0.38135    test-rmse:1.56244
[9]    train-rmse:0.33632    test-rmse:1.57958
[10]   train-rmse:0.28438    test-rmse:1.58056
[11]   train-rmse:0.26198    test-rmse:1.58829
[12]   train-rmse:0.23728    test-rmse:1.60317
[13]   train-rmse:0.21035    test-rmse:1.59406
[14]   train-rmse:0.18668    test-rmse:1.60318
[15]   train-rmse:0.16417    test-rmse:1.59872
```

|      |                    |                   |
|------|--------------------|-------------------|
| [16] | train-rmse:0.15322 | test-rmse:1.60157 |
| [17] | train-rmse:0.13933 | test-rmse:1.60271 |
| [18] | train-rmse:0.12366 | test-rmse:1.60606 |
| [19] | train-rmse:0.11291 | test-rmse:1.60355 |
| [20] | train-rmse:0.10587 | test-rmse:1.60410 |
| [21] | train-rmse:0.10023 | test-rmse:1.60809 |
| [22] | train-rmse:0.09571 | test-rmse:1.60888 |
| [23] | train-rmse:0.08612 | test-rmse:1.61020 |
| [24] | train-rmse:0.08365 | test-rmse:1.61056 |
| [25] | train-rmse:0.08099 | test-rmse:1.61055 |
| [26] | train-rmse:0.06964 | test-rmse:1.60636 |
| [27] | train-rmse:0.06130 | test-rmse:1.60454 |
| [28] | train-rmse:0.05457 | test-rmse:1.60606 |
| [29] | train-rmse:0.05111 | test-rmse:1.60840 |
| [30] | train-rmse:0.04773 | test-rmse:1.60865 |
| [31] | train-rmse:0.04429 | test-rmse:1.60893 |

## Naive Model 1

- **Always predicts** the mean effluent chlorine residual of the **imputed training dataset**

```
y_pred_naive_MBBR = np.full(y_test_MBBR.shape, y_train_MBBR.mean())
```

## Naive Model 2

- **Always predicts** the mean effluent chlorine residual of the **Non-imputed (raw) training dataset**

```
y_pred_naive_orig_MBBR = np.full(y_test_MBBR.shape, y_train_orig_MBBR.mean())
```

## Model Evaluation

```
def compute_metrics(y_pred, y_test):
    std_obs = np.std(y_test)
    std_sim = np.std(y_pred)

    mean_obs = np.mean(y_test)
    mean_sim = np.mean(y_pred)

    # Computing correlation
    r = np.corrcoef(y_test, y_pred)[0, 1]

    # Computing KGE
    alpha = std_sim / std_obs
    beta = mean_sim / mean_obs

    kge = 1 - np.sqrt(np.square(r - 1) + np.square(alpha - 1) + np.square(beta - 1))

    # PBIAS Calculation
    pbias = np.sum((y_test - y_pred)) / np.sum(y_test) * 100

    # Computing NSE
    nse = 1 - (np.sum((y_test - y_pred)**2)) / (np.sum((y_test - np.mean(y_test))**2))

    if nse > 0.35:
        nse = (nse, 'good')
    else:
        nse = (nse, 'bad')
    if abs(pbias) < 15:
        pbias = (abs(pbias), 'good')
    else:
        pbias = (abs(pbias), 'bad')
    if kge > -0.41:
        kge = (kge, 'good')
    else:
        kge = (kge, 'bad')

    return(nse, pbias, kge)

def compute_nrmse(y_true, y_pred):
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
```



```
nrmse = rmse / (np.max(y_true) - np.min(y_true))
return nrmse
```

```
non_imputed_mask_MBBR = ~np.isnan(y_test_orig_MBBR)
```

## ✓ Model Metrics evaluated on Imputed Test Set

### ✓ Optimized XGBoost 1

```
nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR, y_test_MBBR)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")
```

```
rmse = mean_squared_error(y_test_MBBR, y_pred_final_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_final_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Final model metrics:

```
NSE: (0.3604048020224956, 'good'),
PBIAS: (5.258534577934008, 'good'),
KGE: (0.49654117196534, 'good')
```

```
Root Mean Squared Error: 1.3451057796734676
Normalized Root Mean Squared Error: 0.263333159685487
```

### ✓ Untuned XGBoost 1

```
nse_naive, pbias_naive, kge_naive = compute_metrics(y_pred_oob_imputed_MBBR, y_test_MBBR)
print(f"Final model metrics:\n\nNSE: {nse_naive}, \nPBIAS: {pbias_naive}, \nKGE: {kge_naive}")
```

```
rmse = mean_squared_error(y_test_MBBR, y_pred_oob_imputed_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_oob_imputed_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Final model metrics:

```
NSE: (0.17250000700780665, 'bad'),
PBIAS: (6.849152214791969, 'good'),
KGE: (0.481718522004173, 'good')
```

```
Root Mean Squared Error: 1.5299873530816124
Normalized Root Mean Squared Error: 0.2995276728820697
```

### ✓ Naive Model 1

```
rmse = mean_squared_error(y_test_MBBR, y_pred_naive_MBBR, squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_naive_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Root Mean Squared Error: 1.6828674290904115  
Normalized Root Mean Squared Error: 0.3294572100803468

### ✓ Naive Model 2

```
rmse = mean_squared_error(y_test_MBBR, y_pred_naive_orig_MBBR, squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR, y_pred_naive_orig_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

↗ Root Mean Squared Error: 1.6828999269810896  
Normalized Root Mean Squared Error: 0.329463572235922

## Model Metrics evaluated on Non-Imputed (Raw) Test Set

### Optimized XGBoost 1

```
nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR[non_imputed_mask_MBBR], y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")
```

```
rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_final_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_final_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

Final model metrics:

```
NSE: (0.3003406737574903, 'bad'),
PBIAS: (9.138080431028751, 'good'),
KGE: (0.46925790696022485, 'good')

Root Mean Squared Error: 1.3844292516749703
Normalized Root Mean Squared Error: 0.27744073179859124
```

### Optimized XGBoost 2

```
nse_final, pbias_final, kge_final = compute_metrics(y_pred_final_MBBR_dropped, y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_final}, \nPBIAS: {pbias_final}, \nKGE: {kge_final}")
```

```
rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_final_MBBR_dropped, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_final_MBBR_dropped)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

Final model metrics:

```
NSE: (-0.5413266020400924, 'bad'),
PBIAS: (20.85801659458364, 'bad'),
KGE: (0.29327000235137435, 'good')

Root Mean Squared Error: 2.054824760065243
Normalized Root Mean Squared Error: 0.411788529071191
```

### Untuned XGBoost 2

```
nse_naive, pbias_naive, kge_naive = compute_metrics(y_pred_oob_MBBR, y_test_MBBR_dropped)
print(f"Final model metrics:\n\nNSE: {nse_naive}, \nPBIAS: {pbias_naive}, \nKGE: {kge_naive}")
```

```
rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_oob_MBBR, squared=False)
print(f"\nRoot Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_oob_MBBR)
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

Final model metrics:

```
NSE: (0.05503191654014583, 'bad'),
PBIAS: (8.674598570215617, 'good'),
KGE: (0.44942410077765016, 'good')

Root Mean Squared Error: 1.608925827702348
Normalized Root Mean Squared Error: 0.3224300255916529
```

### Naive Model 1

```
rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_naive_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"Root Mean Squared Error: {rmse}")
```

```
nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_naive_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")
```

```

Root Mean Squared Error: 1.6631881676095377
Normalized Root Mean Squared Error: 0.33330424200591935

```

## Naive Model 2

```

rmse = mean_squared_error(y_test_MBBR_dropped, y_pred_naive_orig_MBBR[non_imputed_mask_MBBR], squared=False)
print(f"Root Mean Squared Error: {rmse}")

```

```

nrmse = compute_nrmse(y_test_MBBR_dropped, y_pred_naive_orig_MBBR[non_imputed_mask_MBBR])
print(f"Normalized Root Mean Squared Error: {nrmse}")

```

```

Root Mean Squared Error: 1.6630941882756551
Normalized Root Mean Squared Error: 0.33328540847207516

```

## Feature Importance

```

# Get feature importance
importance_MBBR = final_model_MBBR.get_score(importance_type='gain')

name_dict_MBBR = {
    'Flow.Rate.Influent..m3.d.': 'Flow Rate Influent',
    'BOD.Influent..ppm.': 'BOD Influent',
    'Total.Coliiform.Effluent..MPN.100mL.': 'Total Coliiform Effluent',
    'pH.Pre.chlorination': 'pH Pre-Chlorination',
    'Chlorine.dosage..L.d.': 'Chlorine Dosage',
    'TSS.Pre.chlorination..ppm.': 'TSS Pre-Chlorination',
    'Total.Coliiform.Influent..MPN.100mL.': 'Total Coliiform Influent',
    'Fecal.Coliiform.Influent..MPN.100mL.': 'Fecal Coliiform Influent',
    'BOD.Pre.chlorination..ppm.': 'BOD Pre-Chlorination',
    'Fecal.Coliiform.Effluent..MPN.100mL.': 'Fecal Coliiform Effluent',
    'COD.Influent..ppm.': 'COD Influent',
    'COD.Pre.chlorination..ppm.': 'COD Pre-Chlorination',
}


# For visualization, it is better to convert it to a DataFrame
importance_df_MBBR = pd.DataFrame({
    'Feature': list(importance_MBBR.keys()),
    'Importance': list(importance_MBBR.values())
})

importance_df_MBBR['Feature'] = importance_df_MBBR['Feature'].replace(name_dict_MBBR)

# Sort the DataFrame by importance
importance_df_MBBR = importance_df_MBBR.sort_values(by='Importance', ascending=False)

# Plot feature importance
plt.figure(figsize=(10, 8))
plt.barh(importance_df_MBBR['Feature'], importance_df_MBBR['Importance'], color='skyblue')
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.title("Feature Importance")
plt.gca().invert_yaxis() # To show the highest importance at the top
plt.show()


```



#### ✓ Data Visualization for Model Evaluation

#### ✓ Optimized XGBoost on Imputed Test Dataset

```
# with imputation
plt.scatter(y_test_MBBR, y_pred_final_MBBR);
plt.xlabel('True Value');
plt.ylabel('Predicted Value');
```




```
# Create an x-axis range based on the length of the series/array
x = range(1, len(y_test_MBBR) + 1)
```

```
# Plotting
plt.figure(figsize=(10, 5))
plt.plot(x, y_test_MBBR, label='True', marker='o')
plt.plot(x, y_pred_final_MBBR, label='Predicted', marker='x')


# Adding labels and title
plt.xlabel('Test Datapoint')
plt.ylabel('Effluent Chlorine Residual')
plt.title('Comparison of True and Predicted Values on Imputed Test Data')
plt.legend()

# Show plot
plt.show()
```



#### ✓ Optimized XGBoost on Non-Imputed (Raw) Test Dataset

```
# without imputation
plt.scatter(y_test_orig_MBBR[non_imputed_mask_MBBR], y_pred_final_MBBR[non_imputed_mask_MBBR])
plt.xlabel('True Value')
plt.ylabel('Predicted Value');
```



```
# Create an x-axis range based on the length of the series/array
x = range(1, len(y_test_orig_MBBR[non_imputed_mask_MBBR]) + 1)
```

```
# Plotting
plt.figure(figsize=(10, 5))
plt.plot(x, y_test_orig_MBBR[non_imputed_mask_MBBR], label='True', marker='o')
plt.plot(x, y_pred_final_MBBR[non_imputed_mask_MBBR], label='Predicted', marker='x')

# Adding labels and title
plt.xlabel('Test Datapoint')
plt.ylabel('Effluent Chlorine Residual')
plt.title('Comparison of True and Predicted Values on Raw Test Data')
plt.legend()

# Show plot
plt.show()
```

